



ELF Crafting

Advance Anti-analysis techniques
for the Linux Platform



Whoami

- Security Researcher at Intezer Labs
- Malware Reverse Engineer and Threat Hunter
- CTF player with amn3s1a
- Radare2 Contributor
- Libelfmaster Contributor

 @ulexec




Outline

- Overview to the Linux Threat Landscape
- Advanced ELF crafting techniques
- Q&A

Overview of the Linux Threat Landscape




Linux Malware Distribution (2019)



Intezer - 2019

APTs deploying Linux based Malware




Shadow Brokers

- jackpop
- NOPEN
- SUCTIONCHAR
- SECONDDATE

Vault7

- OutlawCountry
- BothanSpy
- Gyrffalcon



APT28

- Fysbis
- Zebrozy

APT29

- SeaDuke

Turla

- Penguin

Gamaredon

- EvilGnome



Winnti Umbrella

- HiddenWasp
- Azazel Based forks

Challenges on Linux Threat Detection

- Low Visibility
- Low Detection rate



Rise of awareness

- At Intezer we uncovered the following threats on 2019. Most of them were completely undetected before we reported them.
 - **Pacha Group** (Crypto-mining group)
 - **HiddenWasp** (Trojan linked to Winnti umbrella)
 - **QnapCrypt** (Ransomware targeting NAS servers)
 - **EvilGnome** (Trojan linked to Gamaredon Group)
 - **WatchBog Cython** (Crypto-mining botnet)

Rise of awareness

- In recent years we have seen a rise of researchers focused on Linux Malware Hunting promoting awareness on Linux based threats.
 - Cisco Talos
 - Netlab360
 - @liuya0904
 - @_rngm_
 - @zom3y3
 - @JiaYu_521
 - @huiwangeth
 - MalwareMustDie
 - @benkow_
 - @unixfreaxjp
 - Intezer
 - @polarply
 - @ulexec

What's happening next

- Linux malware is doomed to increase in visibility.
- Linux threat detection performance will improve.
- Threat actors will invest more resources to make their implants more evasive and complex

This talk will present some techniques that may start appearing on the wild for defenders awareness as Linux malware evolves in complexity.

Advanced ELF Crafting Techniques



Advanced ELF Crafting

- ELF parsing struggles
 - A word about sections
 - Breaking naive parsers with 1 byte
 - Hiding dynamic entries
- Relocation Hijacking
 - EPO on PIE binaries
 - Hiding constructors in dynamic and static binaries
- Experimental Practices
 - Removing import strings from .dynstr string table

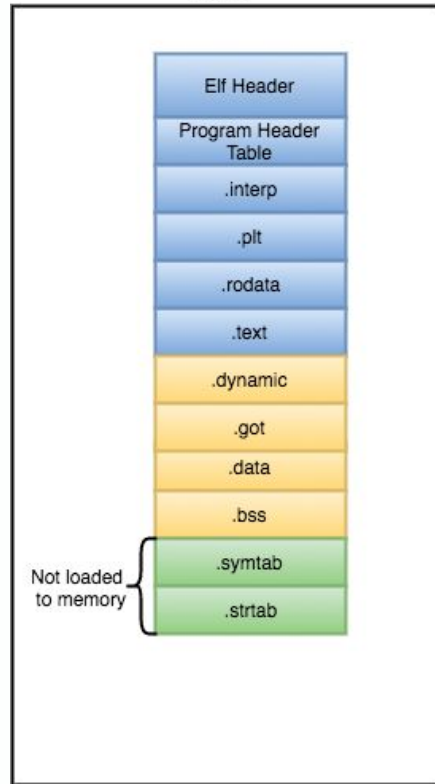
ELF Parsing Struggles




A word about sections

- Sections comprise all information needed for linking a target object file to build a runnable ELF executable.
- Sections are needed on link-time but they are not needed on run-time.
- Compiled ELF binaries will have a Section Header table (an array of *Elf*_Shdr* entries).
- Sections can hold different content ranging from code, strings, relocations or symbols.

Disk



Memory



- Read-Execute
- Read-Write
- Not loadable

A word about sections

- Section information can be completely removed from the binary or modified.
- This can be abused by attackers to misguide analyst or parsers that heavily rely on section information
- Example use-cases:
 - Symbol scrambling
 - Section scrambling

Example outcome

```
ulxec intezer ~ R2CON2019 > section_scamling $ readelf -d ./ls
readelf: Error: Section 0 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 24 for the rest of this dump)
readelf: Error: Section 3 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 16 for the rest of this dump)
readelf: Error: Section 4 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 24 for the rest of this dump)
readelf: Error: Section 11 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 16 for the rest of this dump)
readelf: Error: Section 14 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 24 for the rest of this dump)
readelf: Error: File contains multiple dynamic symbol tables
readelf: Error: File contains multiple dynamic symbol tables
readelf: Error: File contains multiple dynamic symbol tables
readelf: Error: Section 23 has invalid sh_entsize of 0000000000000008
readelf: Error: (Using the expected size of 24 for the rest of this dump)
readelf: Error: Section 24 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 16 for the rest of this dump)
readelf: Error: File contains multiple dynamic symbol tables
readelf: Error: Section 26 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 24 for the rest of this dump)
readelf: Warning: the .dynamic section is not contained within the dynamic segment

Dynamic section at offset 0x3e70 contains 1 entry:
  Tag                Type                Name/Value
0x90660021c15225ff (<unknown>: 90660021c15225ff) 0x90660021c16a25ff
ulxec intezer ~ R2CON2019 > section_scamling $
```

Example outcome

```
ulexec intezer ~ R2CON2019 > section_scramling $ readelf -d ./ls
readelf: Error: Section 0 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 24 for the rest of this dump)
readelf: Error: Section 3 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 16 for the rest of this dump)
readelf: Error: Section 4 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 24 for the rest of this dump)
readelf: Error: Section 11 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 16 for the rest of this dump)
readelf: Error: Section 14 has invalid sh_entsize of 0000000000000000
ulexec intezer ~ R2CON2019 > section_scramling $ objdump -D ./ls
objdump: ./ls: unknown type [0xc] section `
objdump: ./ls: Bad value
readelf: Error: File contains multiple dynamic symbol tables
readelf: Error: Section 23 has invalid sh_entsize of 0000000000000008
readelf: Error: (Using the expected size of 24 for the rest of this dump)
readelf: Error: Section 24 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 16 for the rest of this dump)
readelf: Error: File contains multiple dynamic symbol tables
readelf: Error: Section 26 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 24 for the rest of this dump)
readelf: Warning: the .dynamic section is not contained within the dynamic segment

Dynamic section at offset 0x3e70 contains 1 entry:
   Tag               Type                             Name/Value
0x90660021c15225ff (<unknown>: 90660021c15225ff) 0x90660021c16a25ff
ulexec intezer ~ R2CON2019 > section_scramling $
```

Example outcome

```
ulexec intezer ~ R2CON2019 > section_scamling $ readelf -d ./ls
readelf: Error: Section 0 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 24 for the rest of this dump)
readelf: Error: Section 3 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 16 for the rest of this dump)
readelf: Error: Section 4 has invalid sh_entsize of 0000000000000000
pwndbg: loaded 171 commands. Type pwndbg [filter] for a list.
pwndbg: created $rebase, $ida gdb functions (can be used with print/break)
BFD: /home/ulexec/R2CON2019/section_scamling/ls: unknown type [0xc] section `
"/home/ulexec/R2CON2019/section_scamling/ls": not in executable format: Bad value
pwndbg> r
Starting program:
No executable file specified.
Use the "file" or "exec-file" command.
pwndbg> 
readelf: Error: (Using the expected size of 16 for the rest of this dump)
readelf: Error: File contains multiple dynamic symbol tables
readelf: Error: Section 26 has invalid sh_entsize of 0000000000000000
readelf: Error: (Using the expected size of 24 for the rest of this dump)
readelf: Warning: the .dynamic section is not contained within the dynamic segment

Dynamic section at offset 0x3e70 contains 1 entry:
  Tag               Type                             Name/Value
0x90660021c15225ff (<unknown>: 90660021c15225ff) 0x90660021c16a25ff
ulexec intezer ~ R2CON2019 > section_scamling $
```

Circumvention

- If sections are crafted is better to neglect them
- Sections can be neglected by zeroing out the following fields in the ELF header
 - shoff
 - shnum
 - shstrndx

```
ulexec > intezer ~ R2CON2019 > section_scramling $ r2 -nn -w ../cp
[0x00000000]> pfo elf64
[0x00000000]> pf elf_header
  ident : 0x00000000 = "\x7fELF\x02\x01\x01"
  type : 0x00000010 = type (enum elf_type) = 0x3 ; ET_DYN
  machine : 0x00000012 = machine (enum elf_machine) = 0x3e ; EM_AMD64
  version : 0x00000014 = 0x00000001
  entry : 0x00000018 = (qword)0x0000000000004640
  phoff : 0x00000020 = (qword)0x0000000000000040
  shoff : 0x00000028 = (qword)0x00000000000221d8
  flags : 0x00000030 = 0x00000000
  ehsize : 0x00000034 = 0x0040
  phentsize : 0x00000036 = 0x0038
  phnum : 0x00000038 = 0x0009
  shentsize : 0x0000003a = 0x0040
  shnum : 0x0000003c = 0x001c
  shstrndx : 0x0000003e = 0x001b
[0x00000000]> .pf.elf_header.shnum=0
[0x00000000]> .pf.elf_header.shoff=0
[0x00000000]> .pf.elf_header.shstrndx=0
[0x00000000]>
```

Circumvention

- Import resolution can be done via .dynstr/.dynsym from PT_DYNAMIC segment entries DT_STRTAB and DT_SYMTAB accordingly

```
Dynamic section at offset 0x1fa38 contains 28 entries:
  Tag          Type          Name/Value
0x0000000000000001 (NEEDED)      Shared library: [libselinux.so.1]
0x0000000000000001 (NEEDED)      Shared library: [libc.so.6]
0x000000000000000c (INIT)        0x3758
0x000000000000000d (FINI)        0x1636c
0x0000000000000019 (INIT_ARRAY)    0x21eff0
0x000000000000001b (INIT_ARRAYSZ) 8 (bytes)
0x000000000000001a (FINI_ARRAY)    0x21eff8
0x000000000000001c (FINI_ARRAYSZ) 8 (bytes)
0x0000000006ffffe5 (GNU_HASH)     0x298
0x0000000000000005 (STRTAB)       0x1180
0x0000000000000006 (SYMTAB)       0x388
0x000000000000000a (STRSZ)        1008 (bytes)
0x000000000000000b (SYMMENT)      24 (bytes)
0x0000000000000015 (DEBUG)        0x0
0x0000000000000003 (PLTGOT)       0x21fc38
```

- Export resolution can be done parsing entries from PT_GNU_EH_FRAME segment
 - https://github.com/intezer/scripts/blob/master/eh_frame_parser.py

1 byte parser breaker

- ELF section manipulation can provoke parsing problems.
- ELF fields can be crafted to incorrectly parse a given ELF image.
- Sometimes not much effort is needed to achieve this outcome.

1 Byte Parser Breaker

```
ulexec ~ R2CON2019 > parser_breaker $ r2 -w -nn ./ls
[0x00000000]> pfo elf64
[0x00000000]> pf elf_header
  ident : 0x00000000 = "\x7fELF\x02\x01\x01"
  type  : 0x00000010 = type (enum elf_type) = 0x3 ; ET_DYN
  machine : 0x00000012 = machine (enum elf_machine) = 0x3e ; EM_AMD64
  version : 0x00000014 = 0x00000001
  entry  : 0x00000018 = (qword)0x00000000000005850
  phoff  : 0x00000020 = (qword)0x0000000000000040
  shoff  : 0x00000028 = (qword)0x00000000000203a0
  flags  : 0x00000030 = 0x00000000
  ehsize : 0x00000034 = 0x0040
  phentsize : 0x00000036 = 0x0038
  phnum  : 0x00000038 = 0x0009
  shentsize : 0x0000003a = 0x0040
  shnum  : 0x0000003c = 0x001c
  shstrndx : 0x0000003e = 0x001b
```

1 Byte Parser Breaker

```
ulexec ~ R2CON2019 > parser_breaker $ r2 -w -nn ./ls
[0x00000000]> pfo elf64
[0x00000000]> pf elf_header
  ident : 0x00000000 = "\x7fELF\x02\x01\x01"
  type  : 0x00000010 = type (enum elf_type) = 0x3 ; ET_DYN
  machine : 0x00000012 = machine (enum elf_machine) = 0x3e ; EM_AMD64
  version : 0x00000014 = 0x00000001
  entry   : 0x00000018 = (qword)0x00000000000005850
  phoff   : 0x00000020 = (qword)0x0000000000000040
  shoff   : 0x00000028 = (qword)0x00000000000203a0
  flags   : 0x00000030 = 0x00000000
  ehsize  : 0x00000034 = 0x0040
  phentsize : 0x00000036 = 0x0038
  phnum    : 0x00000038 = 0x0009
  shentsize : 0x0000003a = 0x0040
  shnum     : 0x0000003c = 0x001c
  shstrndx  : 0x0000003e = 0x001b
```


1 Byte Parser Breaker

Array of 16 bytes containing identification flags about the file, which serve to decode and interpret the file's contents. Examples of these identification flags include:

- EI_MAG0-3: ELF magic
- EI_CLASS: File class.
- **EI_DATA: File's data encoding.**
- EI_VERSION: File's version.
- EI_OSABI: OS/ABI identification.
- EI_ABIVERSION: ABI version
- EI_PAD: Start of padding bytes.
- EI_NIDENT: Size of ei_ident.

1 Byte Parser Breaker

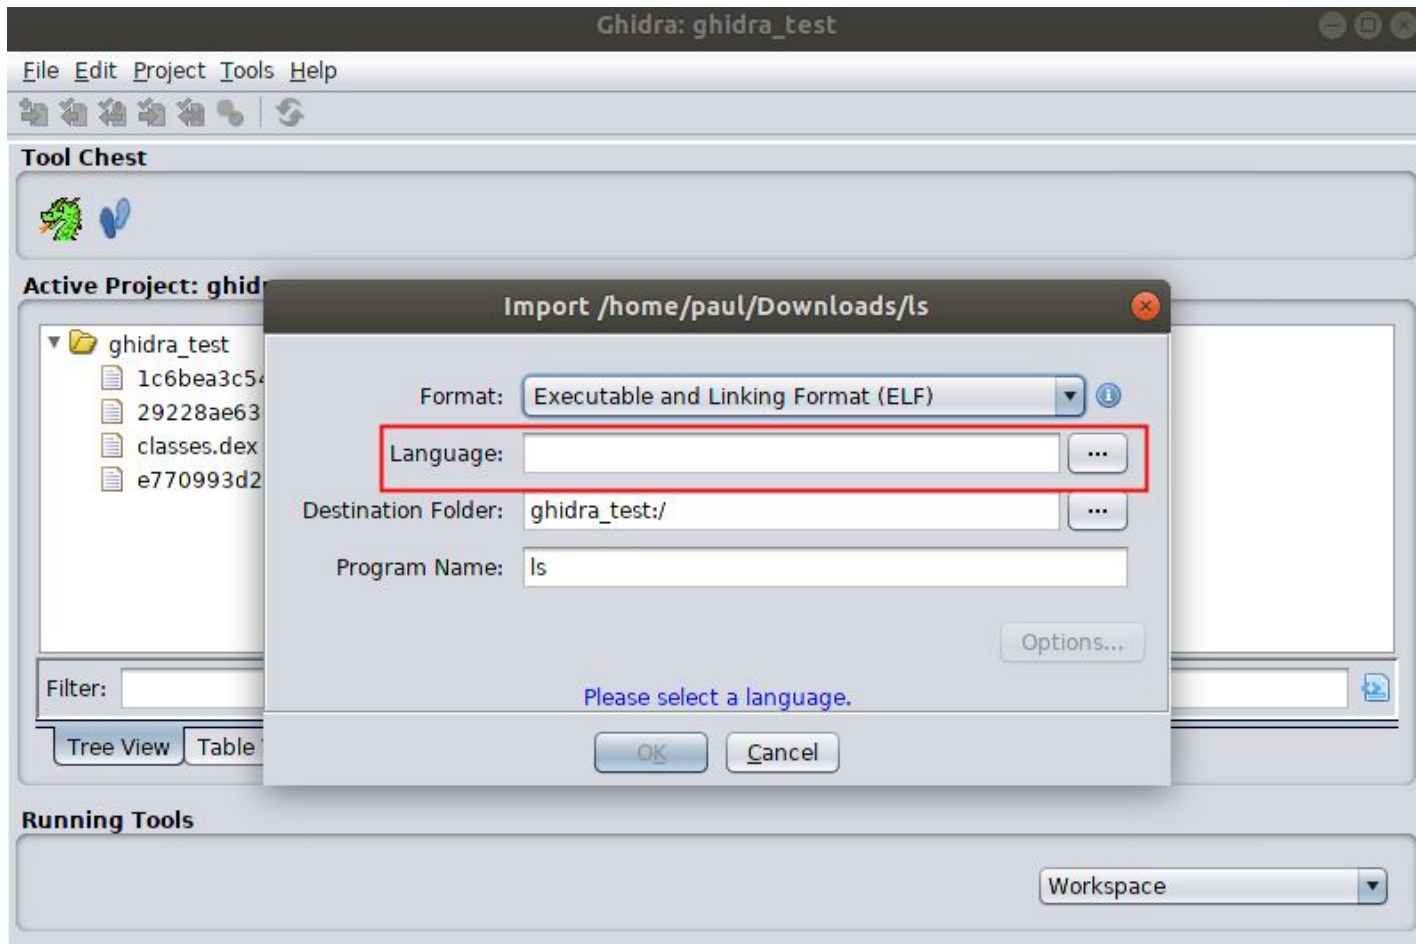
```
[0x00000000]> pf elf_header
ident :
    struct<elf_ident>
    magic : 0x00000000 = "\x7fELF"
    class : 0x00000004 = class (enum elf_class) = 0x2 ; ELFCLASS64
    data : 0x00000005 = data (enum elf_data) = 0x2 ; ELFDATA2MSB
    version : 0x00000006 = version (enum elf_hdr_version) = 0x1 ; EV_CURRENT
    type : 0x00000010 = type (enum elf_type) = 0x3 ; ET_DYN
    machine : 0x00000012 = machine (enum elf_machine) = 0x3e ; EM_AMD64
    version : 0x00000014 = version (enum elf_obj_version) = 0x1 ; EV_CURRENT
    entry : 0x00000018 = 0x00005850
    phoff : 0x0000001c = 0x00000000
    shoff : 0x00000020 = 0x00000040
    flags : 0x00000024 = 0x00000000
    ehsize : 0x00000028 = 0x03a0
    phentsize : 0x0000002a = 0x0002
    phnum : 0x0000002c = 0x0000
    shentsize : 0x0000002e = 0x0000
    shnum : 0x00000030 = 0x0000
    shstrndx : 0x00000032 = 0x0000
```

1 Byte

```
ulexec intezer ~ R2CON2019 > parser_breaker $ readelf -h ./ls
ELF Header:
  Magic:   7f 45 4c 46 02 02 01 00 00 00 00 00 00 00 00 00
  Class:                               ELF64
  Data:                               2's complement, big endian
  Version:                             1 (current)
  OS/ABI:                              UNIX - System V
  ABI Version:                         0
  Type:                                <unknown>: 300
  Machine:                             <unknown>: 0x3e00
  Version:                             0x1000000
  Entry point address:                 0x5058000000000000
  Start of program headers:            4611686018427387904 (bytes into file)
  Start of section headers:            -6916682403687694336 (bytes into file)
  Flags:                                0x0
  Size of this header:                  16384 (bytes)
  Size of program headers:              14336 (bytes)
  Number of program headers:            2304
  Size of section headers:              16384 (bytes)
  Number of section headers:            7168
  Section header string table index:    6912
readelf: Warning: The e_shentsize field in the ELF header is larger than the size of an ELF section
header
readelf: Error: Reading 117440512 bytes extends past end of file for section headers
readelf: Warning: The e_phentsize field in the ELF header is larger than the size of an ELF program
header
readelf: Error: Reading 33030144 bytes extends past end of file for program headers
ulexec intezer ~ R2CON2019 > parser_breaker $ readelf -l ./ls
readelf: Warning: The e_shentsize field in the ELF header is larger than the size of an ELF section
header
readelf: Error: Reading 117440512 bytes extends past end of file for section headers

Elf file type is <unknown>: 300
Entry point 0x5058000000000000
There are 2304 program headers, starting at offset 4611686018427387904
readelf: Warning: The e_phentsize field in the ELF header is larger than the size of an ELF program
header
readelf: Error: Reading 33030144 bytes extends past end of file for program headers
ulexec intezer ~ R2CON2019 > parser_breaker $
```







ERROR LOADING SESSION

There was an error loading this session! If this is the first time you're trying to view this session, please check that you uploaded a valid binary or BNDB.

```

ulexec  intezer  ~  R2CON2019 > parser_breaker  $  r2 ./ls
[0x5058000000000000]> pd 10
    |-- entry0:
    |-- rip:
    0x5058000000000000      ff      invalid
    0x5058000000000001      ff      invalid
    0x5058000000000002      ff      invalid
    0x5058000000000003      ff      invalid
    0x5058000000000004      ff      invalid
    0x5058000000000005      ff      invalid
    0x5058000000000006      ff      invalid
    0x5058000000000007      ff      invalid
    0x5058000000000008      ff      invalid
    0x5058000000000009      ff      invalid
[0x5058000000000000]> ii
[Imports]
Num  Vaddr      Bind      Type Name

[0x5058000000000000]> S
|ERROR| Invalid command 'S' (0x53)
[0x5058000000000000]> is
[Sections]
Nm Paddr      Size Vaddr      Memsz Perms Name
00 0x00000000 133792 0x00010000 133792 -rwx uphdr

```

1 Byte Parser Breaker

- Most static parsers attempt to find the endianness of the file before the machine type.
- The e_machine field will be interpreted according to the endianness field.
- The only disassembler that I've seen that handled this anomaly gracefully was IDA.
- This technique could impact **dynamic analysis automation** and **sandbox performance** of the subject ELF file.

1 Byte Parser Breaker

The reason this technique can be enforced without interacting with the kernel's ELF loader is because the endianness is set on compile time in the linux kernel as the CPU endianness.

```
arch > x86 > include > asm > C elf.h > ELF_DATA
20 typedef struct user_i387_struct elf_fpregset_t;
21
22 #ifdef __i386__
23
24 typedef struct user_fxsr_struct elf_fpxregset_t;
25
26 #define R_386_NONE 0
27 #define R_386_32 1
28 #define R_386_PC32 2
29 #define R_386_GOT32 3
30 #define R_386_PLT32 4
31 #define R_386_COPY 5
32 #define R_386_GLOB_DAT 6
33 #define R_386_JMP_SLOT 7
34 #define R_386_RELATIVE 8
35 #define R_386_GOTOFF 9
36 #define R_386_GOTPC 10
37 #define R_386_NUM 11
38
39 /*
40  * These are used to set parameters in the core dumps.
41  */
42 #define ELF_CLASS ELFCLASS32
43 #define ELF_DATA ELFDATA2LSB
44 #define ELF_ARCH EM_386
```


Circumvention

If an ELF file seems to have an unknown architecture and readelf's output seems highly broken,

Restore endianness byte accordingly and attempt further parsing heuristics.

```
ulexec intezer ~ R2CON2019 > parser_breaker $ r2 -w -nn ./ls
[0x00000000]> wx 0x01 @+5
[0x00000000]> pfo elf32
[0x00000000]> pf elf_header
ident :
    struct<elf_ident>
    magic : 0x00000000 = "\x7fELF"
    class : 0x00000004 = class (enum elf_class) = 0x2 ; ELFCLASS64
    data : 0x00000005 = data (enum elf_data) = 0x1 ; ELFDATA2LSB
    version : 0x00000006 = version (enum elf_hdr_version) = 0x1 ; EV_CURRENT
    type : 0x00000010 = type (enum elf_type) = 0x3 ; ET_DYN
    machine : 0x00000012 = machine (enum elf_machine) = 0x3e ; EM_AMD64
    version : 0x00000014 = version (enum elf_obj_version) = 0x1 ; EV_CURRENT
    entry : 0x00000018 = 0x00005850
    phoff : 0x0000001c = 0x00000000
    shoff : 0x00000020 = 0x00000040
    flags : 0x00000024 = 0x00000000
    ehsize : 0x00000028 = 0x03a0
    phentsize : 0x0000002a = 0x0002
    phnum : 0x0000002c = 0x0000
    shentsize : 0x0000002e = 0x0000
    shnum : 0x00000030 = 0x0000
    shstrndx : 0x00000032 = 0x0000
```

Hiding dynamic entries

- Dynamically compiled ELF files will always contain a segment of type PT_DYNAMIC
- This segment contains all needed information needed for the dynamic linking \ to take place.
- Dynamically linked binaries with a defectuous dynamic segment will crash

Hiding dynamic entries

- Android packer scene is much more mature in regards to evasion techniques than standalone ELF malware is.
- Most ELFs found in Android applications tend to be shared objects.
- Successfully analysis of shared objects heavily relies on visibility of dynamic linking artifacts. (init_array, exports, ... etc)
- The following technique implements a way to hide dynamic entries from ELF shared objects from the PT_DYNAMIC segment

Detecting anomalies

```
ulexec  intezer  ~  R2CON2019 > hidden_dynamic_entries $ r2 libSDKRelativeJNI.so
[0x00000000]> aa
[x] Analyze all flags starting with sym. and entry0 (aa)
[0x00000000]> is
[Symbols]
Num Paddr      Vaddr      Bind      Type Size Name
No symbols
[0x00000000]> !readelf -d ./libSDKRelativeJNI.so
readelf: Error: no .dynamic section in the dynamic segment
Section parsing error
Dynamic section at offset 0x41464 contains 11 entries:
  Tag      Type      Name/Value
0x00000001 (NEEDED)  0x8328
0x00000001 (NEEDED)  0x8337
0x00000001 (NEEDED)  0x8347
0x00000001 (NEEDED)  0x8351
0x00000001 (NEEDED)  0x835e
0x00000001 (NEEDED)  0x8366
0x00000001 (NEEDED)  0x8373
0x00000001 (NEEDED)  0x837b
0x00000001 (NEEDED)  0x8383
0x0000000e (SONAME)  0x838c
0x00000000 (NULL)  0x0
[0x00000000]> 
```

Trying to circumvent anomalies

```
ulxex  intezer  ~ R2CON2019 > hidden_dynamic_entries $ r2 -w -nn libSDKRelativeJNI_.so
[0x00000000]> pfo elf32
[0x00000000]> pf elf_header
    ident :
        struct<elf_ident>
        magic : 0x00000000 = "\x7fELF"
        class : 0x00000004 = class (enum elf_class) = 0x1 ; ELFCLASS32
        data : 0x00000005 = data (enum elf_data) = 0x1 ; ELFDATA2LSB
        version : 0x00000006 = version (enum elf_hdr_version) = 0x1 ; EV_CURRENT
        type : 0x00000010 = type (enum elf_type) = 0x3 ; ET_DYN
        machine : 0x00000012 = machine (enum elf_machine) = 0x28 ; EM_ARM
        version : 0x00000014 = version (enum elf_obj_version) = 0x1 ; EV_CURRENT
        entry : 0x00000018 = 0x00000000
        phoff : 0x0000001c = 0x00000034
        shoff : 0x00000020 = 0x000414e4
        flags : 0x00000024 = 0x05000000
        ehsize : 0x00000028 = 0x0034
        phentsize : 0x0000002a = 0x0020
        phnum : 0x0000002c = 0x0004
        shentsize : 0x0000002e = 0x0028
        shnum : 0x00000030 = 0x0004
        shstrndx : 0x00000032 = 0x0001
[0x00000000]> .pf.elf_header.shnum=0
[0x00000000]> .pf.elf_header.shoff=0
[0x00000000]> .pf.elf_header.shstrndx=0
```

More anomalies

```
ulexec > intezer > ~ > R2CON2019 > hidden_dynamic_entries $ readelf -h ./libSDKRelativeJNI.so

ELF Header:
  Magic:   7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
  Class:             ELF32
  Data:              2's complement, little endian
  Version:           1 (current)
  OS/ABI:            UNIX - System V
  ABI Version:       0
  Type:              DYN (Shared object file)
  Machine:           ARM
  Version:           0x1
  Entry point address: 0x0
  Start of program headers: 52 (bytes into file)
  Start of section headers: 0 (bytes into file)
  Flags:             0x5000000, Version5 EABI
  Size of this header: 52 (bytes)
  Size of program headers: 32 (bytes)
  Number of program headers: 4
  Size of section headers: 40 (bytes)
  Number of section headers: 0
  Section header string table index: 0
readelf: Error: the dynamic segment offset + size exceeds the size of the file
ulexec > intezer > ~ > R2CON2019 > hidden_dynamic_entries $ readelf -d ./libSDKRelativeJNI.so

readelf: Error: the dynamic segment offset + size exceeds the size of the file

There is no dynamic section in this file.
ulexec > intezer > ~ > R2CON2019 > hidden_dynamic_entries $ python -c 'print "/x00" * 100' >>
./libSDKRelativeJNI.so
```

No profit

```
ulxec > intezer ~ > R2CON2019 > hidden_dynamic_entries $ > readelf -d ./libSDKRelativeJNI_.so
```

Dynamic section at offset 0x41464 contains 11 entries:

Tag	Type	Name/Value
0x00000001	(NEEDED)	0x8328
0x00000001	(NEEDED)	0x8337
0x00000001	(NEEDED)	0x8347
0x00000001	(NEEDED)	0x8351
0x00000001	(NEEDED)	0x835e
0x00000001	(NEEDED)	0x8366
0x00000001	(NEEDED)	0x8373
0x00000001	(NEEDED)	0x837b
0x00000001	(NEEDED)	0x8383
0x0000000e	(SONAME)	0x838c
0x00000000	(NULL)	0x0

Focusing on Program Headers

```
ulexec > intezer > ~ > R2CON2019 > hidden_dynamic_entries $ readelf -l ./libSDKRelativeJNI_.so
```

Elf file type is DYN (Shared object file)
Entry point 0x0
There are 4 program headers, starting at offset 52

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
LOAD	0x000000	0x00000000	0x00000000	0x3db60	0x7e91c	R E	0x1000
LOAD	0x03e688	0x00080688	0x00080688	0x02dcc	0x08c00	RW	0x1000
DYNAMIC	0x041464	0x00082ac8	0x00082ac8	0x00120	0x00120	RW	0x4
EXIDX	0x07504c	0x0007504c	0x0007504c	0x03968	0x03968	R	0x4

Focusing on Program Headers

```
ulexec > intezer > ~ > R2CON2019 > hidden_dynamic_entries $ readelf -l ./libSDKRelativeJNI_.so
```

Elf file type is DYN (Shared object file)
Entry point 0x0
There are 4 program headers, starting at offset 52

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
LOAD	0x000000	0x00000000	0x00000000	0x3db60	0x7e91c	R E	0x1000
LOAD	0x03e688	0x00080688	0x00080688	0x02dcc	0x08c00	RW	0x1000
DYNAMIC	0x041464	0x00082ac8	0x00082ac8	0x00120	0x00120	RW	0x4
EXIDX	0x07504c	0x0007504c	0x0007504c	0x03968	0x03968	R	0x4

Inconsistencies between memory and disk layouts

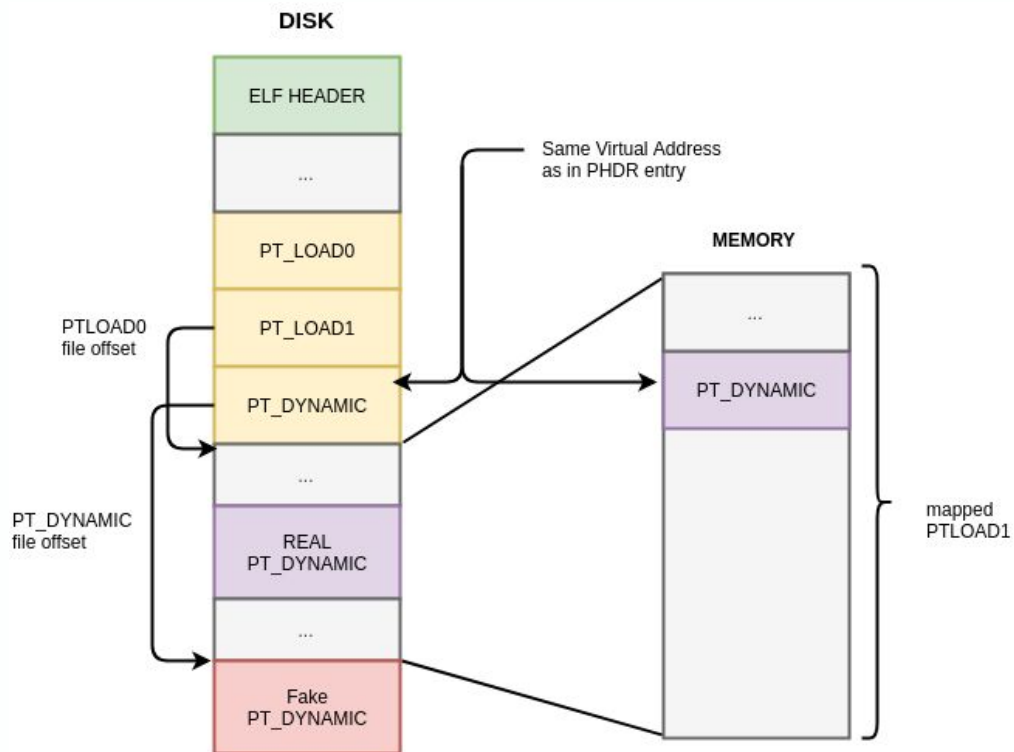
```
[0x00082ac8 [Xadvc] 95% 1848 ./libSDKRelativeJNI_.so]> xc @ segment.DYNAMIC
- offset -   0 1 2 3 4 5 6 7 8 9 A B C D E F 0123456789ABCDEF comment
0x00082ac8 0300 0000 982e 0800 0200 0000 b802 0000 ..... ; segment.DYNAMIC
0x00082ad8 1700 0000 1c11 0100 1400 0000 1100 0000 .....
0x00082ae8 1100 0000 fcd5 0000 1200 0000 203b 0000 ..... ;..
0x00082af8 1300 0000 0800 0000 faff ff6f 6007 0000 ..... o`..
0x00082b08 0600 0000 4801 0000 0b00 0000 1000 0000 ...H.....
0x00082b18 0500 0000 983b 0000 0a00 0000 a183 0000 ....;.....
0x00082b28 0400 0000 3cbf 0000 0100 0000 2883 0000 ....<.....(..
0x00082b38 0100 0000 3783 0000 0100 0000 4783 0000 ....7.....G..
0x00082b48 0100 0000 5183 0000 0100 0000 5e83 0000 ....Q.....^...
0x00082b58 0100 0000 6683 0000 0100 0000 7383 0000 ....f.....s...
0x00082b68 0100 0000 7b83 0000 0100 0000 8383 0000 ....{.....

[0x00041464 [Xadvc] 0% 1848 ./libSDKRelativeJNI_.so]> xc @ elf_phdr+267312 # 0x41464
- offset -   0 1 2 3 4 5 6 7 8 9 A B C D E F 0123456789ABCDEF comment
0x00041464 0100 0000 2883 0000 0100 0000 3783 0000 ....(.....7...
0x00041474 0100 0000 4783 0000 0100 0000 5183 0000 ....G.....Q...
0x00041484 0100 0000 5e83 0000 0100 0000 6683 0000 ....^.....f...
0x00041494 0100 0000 7383 0000 0100 0000 7b83 0000 ....s.....{...
0x000414a4 0100 0000 8383 0000 0e00 0000 8c83 0000 .....
0x000414b4 0000 0000 0000 0000 0000 0000 0000 0000 .....
0x000414c4 0000 0000 0000 0000 0000 0000 0000 0000 .....
0x000414d4 0000 0000 0000 0000 0000 0000 0000 0000 .....
0x000414e4 0000 0000 0000 0000 0000 0000 0000 0000 .....
0x000414f4 0000 0000 0000 0000 0000 0000 0000 0000 .....
```

Memory

Disk

Dynamic segment obfuscation technique



Circumvention

- This technique can be detected checking if the following heuristic does not match:

$$\begin{aligned} \text{dynamic segment offset} &= \text{data segment offset} \\ &+ (\text{vaddr dynamic segment} - \text{vaddr data segment}) \end{aligned}$$

Circumvention

```
ulexec  intezer  ~  R2CON2019 > hidden_dynamic_entries $ r2 -nn -w libSDKRelativeJNI.so
[0x00000000]> !readelf -l ./libSDKRelativeJNI.so

Elf file type is DYN (Shared object file)
Entry point 0x0
There are 4 program headers, starting at offset 52

Program Headers:
  Type           Offset   VirtAddr   PhysAddr   FileSiz MemSiz  Flg Align
  LOAD           0x000000 0x00000000 0x00000000 0x3db60 0x7e91c R E 0x1000
  LOAD           0x03e688 0x00080688 0x00080688 0x02dcc 0x08c00 RW 0x1000
  DYNAMIC        0x041464 0x00082ac8 0x00082ac8 0x00120 0x00120 RW 0x4
  EXIDX          0x07504c 0x0007504c 0x0007504c 0x03968 0x03968 R 0x4
[0x00000000]> ?v 0x03e688 + (0x00082ac8 - 0x00080688)
0x40ac8
[0x00000000]> pfo elf32
[0x00000000]> pf.elf_header.phoff
  phoff : 0x0000001c = 0x00000034
[0x00000000]> pf.elf_header.phentsize
  phentsize : 0x0000002a = 0x0020
[0x00000000]> s 0x34 + (0x20 * 2)
[0x00000074]> pf.elf_phdr
  type : 0x00000074 = type (enum elf_p_type) = 0x2 ; PT_DYNAMIC
  offset : 0x00000078 = 0x00041464
  vaddr : 0x0000007c = 0x00082ac8
  paddr : 0x00000080 = 0x00082ac8
  filesz : 0x00000084 = 0x00000120
  memsz : 0x00000088 = 0x00000120
  flags : 0x0000008c = flags (enum elf_p_flags) = 0x6 ; PF_Read_Write
  align : 0x00000090 = 0x00000004
[0x00000074]> .pf.elf_phdr.offset=0x40ac8
```

Profit

```
Dynamic section at offset 0x40ac8 contains 32 entries:
Type      Type      Name/Value
0x00000003 (PLTGOT)      0x82e98
0x00000002 (PLTRELSZ)    696 (bytes)
0x00000017 (JMPREL)      0x1111c
0x00000014 (PLTREL)      REL
0x00000011 (REL)         0xd5fc
0x00000012 (RELSZ)       15136 (bytes)
0x00000013 (RELENT)      8 (bytes)
0x6ffffffa (RELCOUNT)    1888
0x00000006 (SYMTAB)      0x148
0x0000000b (SYMENT)      16 (bytes)
0x00000005 (STRTAB)      0x3b98
0x0000000a (STRSZ)       33697 (bytes)
0x00000004 (HASH)        0xbf3c
0x00000001 (NEEDED)      Shared library: [libdjiivideo.so]
0x00000001 (NEEDED)      Shared library: [libGLv1_CM.so]
0x00000001 (NEEDED)      Shared library: [liblog.so]
0x00000001 (NEEDED)      Shared library: [libGLv2.so]
0x00000001 (NEEDED)      Shared library: [libz.so]
0x00000001 (NEEDED)      Shared library: [libstdc++.so]
0x00000001 (NEEDED)      Shared library: [libm.so]
0x00000001 (NEEDED)      Shared library: [libc.so]
0x00000001 (NEEDED)      Shared library: [libdl.so]
0x0000000e (SONAME)      Library soname: [libSDKRelativeJNI.so]
0x0000001a (FINI_ARRAY)  0x80de0
0x0000001c (FINI_ARRAYSZ) 8 (bytes)
0x00000019 (INIT_ARRAY)  0x80de8
0x0000001b (INIT_ARRAYSZ) 32 (bytes)
0x00000010 (SYMBOLIC)    0x0
0x0000001e (FLAGS)       SYMBOLIC BIND_NOW
0x6ffffffb (FLAGS_1)     Flags: NOW
0x0000000c (INIT)        0x113d8
0x00000000 (NULL)        0x0
```


Profit

```
readelf: Error: no .dynamic section in the dynamic segment
```

```
Dynamic section at offset 0x40ac8 contains 32 entries:
```

Tag	Type	Name/Value
0x00000003	(PLTGOT)	0x82e98
0x00000002	(PLTRELSZ)	696 (bytes)
0x00000017	(JMPREL)	0x1111c
0x00000014	(PLTREL)	REL
0x00000011	(REL)	0xd5fc
0x00000012	(RELSZ)	15136 (bytes)
0x00000013	(RELENT)	8 (bytes)
0x6ffffffa	(RELCOUNT)	1888
0x00000006	(SYMTAB)	0x148
0x0000000b	(SYMENT)	16 (bytes)
0x00000005	(STRTAB)	0x3b98
0x0000000a	(STRSZ)	33607 (bytes)

```
ulxexec > intezer ~ R2CON2019 > hidden_dynamic_entries 1 $ r2 libSDKRelativeJNI.so
```

```
[0x00000000]> is~?
```

```
05
```

0x00000001	(NEEDED)	Shared library: [libOLE32V2.so]
0x00000001	(NEEDED)	Shared library: [libz.so]
0x00000001	(NEEDED)	Shared library: [libstdc++.so]
0x00000001	(NEEDED)	Shared library: [libm.so]
0x00000001	(NEEDED)	Shared library: [libc.so]
0x00000001	(NEEDED)	Shared library: [libdl.so]
0x0000000e	(SONAME)	Library soname: [libSDKRelativeJNI.so]
0x0000001a	(FINI_ARRAY)	0x80de0
0x0000001c	(FINI_ARRAYSZ)	8 (bytes)
0x00000019	(INIT_ARRAY)	0x80de8
0x0000001b	(INIT_ARRAYSZ)	32 (bytes)
0x00000010	(SYMBOLIC)	0x0
0x0000001e	(FLAGS)	SYMBOLIC BIND_NOW
0x6ffffffb	(FLAGS_1)	Flags: NOW
0x0000000c	(INIT)	0x113d8
0x00000000	(NULL)	0x0



Relocation Hijacking



Relocation Hijacking

- ELF files contain a very sophisticated relocation system.
- Relocation information is held in entries located in specific relocation sections within an ELF object.
- These entries are implemented in the form of structures. There are two different Relocation entry structures: `Elfxx_Rel` and `Elfxx_Rela`:

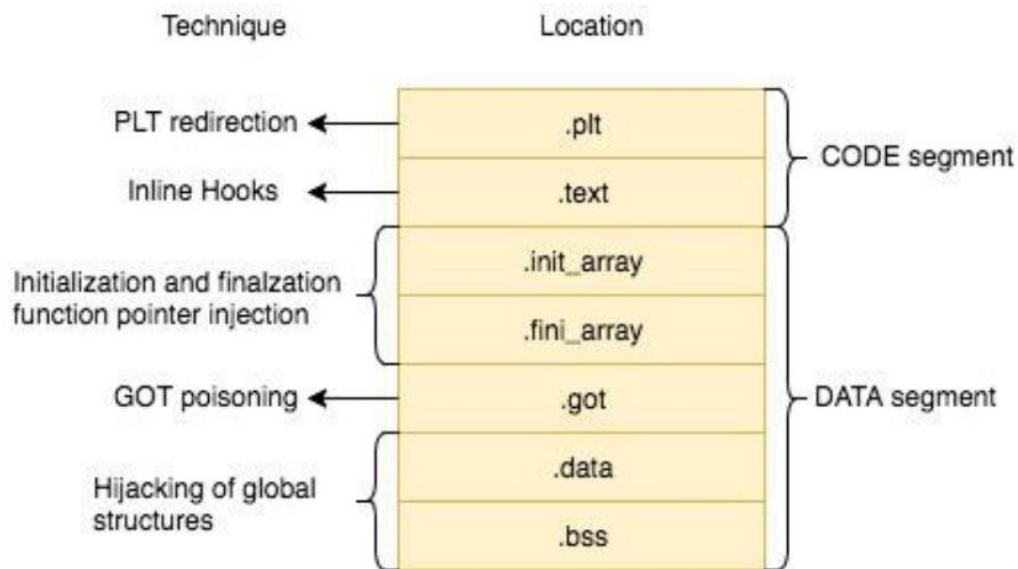
Relocation Hijacking

```
typedef struct {  
    Elf32_Addr    r_offset;  
    Elf32_Word    r_info;  
} Elf32_Rel;  
  
typedef struct {  
    Elf32_Addr    r_offset;  
    Elf32_Word    r_info;  
    Elf32_Sword    r_addend;  
} Elf32_Rela;
```

Relocation Hijacking

- Code injection techniques for the ELF file format exists to build custom tools or craft more complex implants.
- However there is still a challenge of pivoting control of execution into the injected payload.
- There are well known techniques abused in the past in order to do this.

Relocation Hijacking



Relocation Hijacking

- These techniques are known for a long time and easily detectable
- However there are ways to hijack relocations in order to pivot control flow in more stealthy ways.
- Two ideas:
 - EPO based on relocation tampering for PIE binaries
 - Hiding constructors in Dynamically and Statically linked binaries

Challenges on malicious relocations

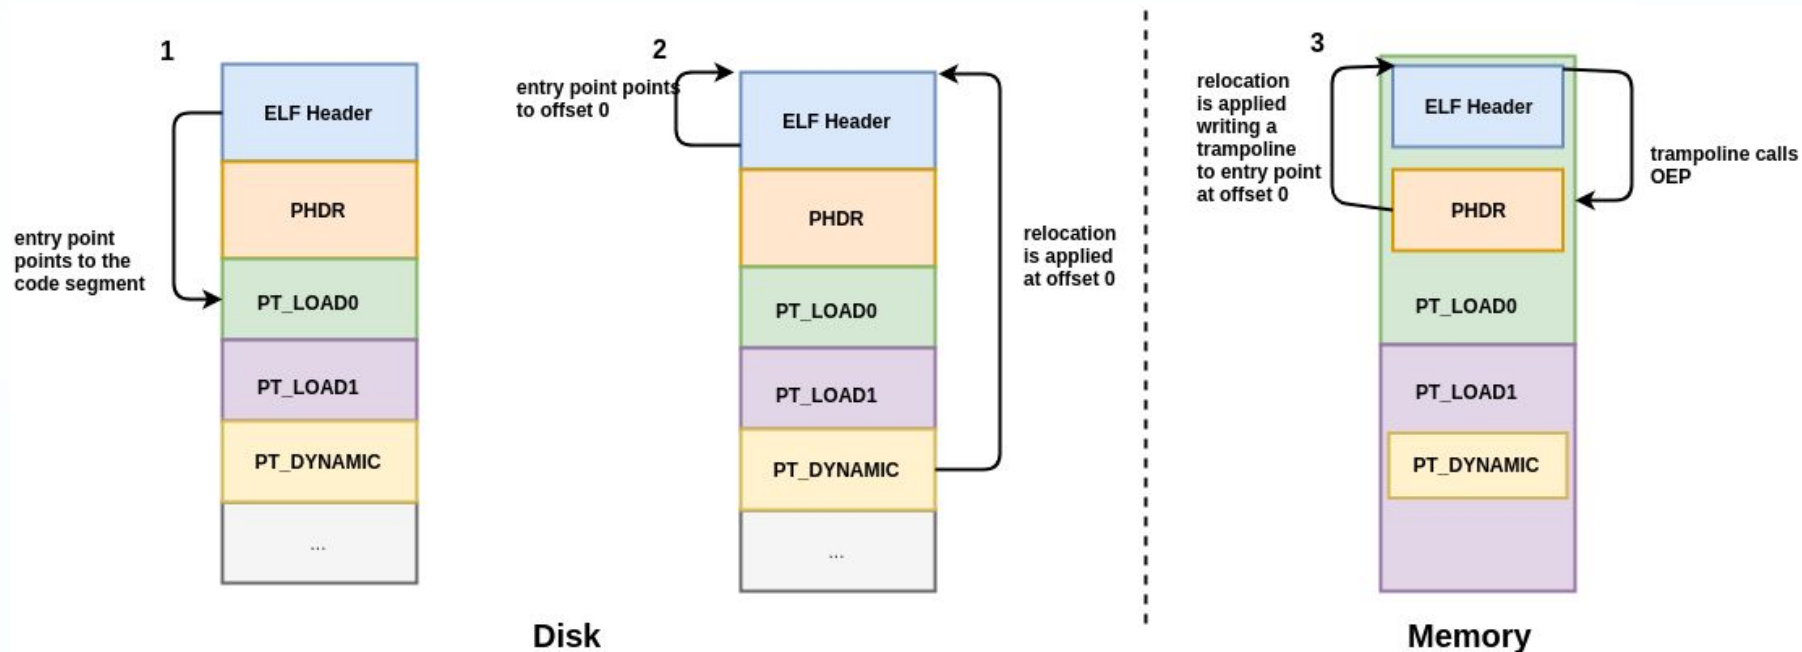
- Relocations essentially enforce hot-patching the binary on load-time.
- Relocation analysis is hard.
- Advanced heuristics have to be applied in order to recognise if a relocation has malicious intentions.
- Relocations are easy to tamper.

EPO on PIE executables

- PIE executables are ET_DYN binaries
- Image base will be chosen on load-time
- Virtual addresses in disk are relative to the chosen image base
- In disk the file's image base is 0.

Demo

Overview



Overview

```
for (i = 0; i < relasz; i++, rela++) {
    Elf64_Sym *sym = &dynsym[ELF64_R_SYM(rela->r_info)];
    if (!strcmp(&dynstr[sym->st_name], symbol_name)) {
        rela->r_offset = 0;
        rela->r_info = (rela->r_info & ~0xf) | R_X86_64_64;
        sym->st_value = 123456;
        rela->r_addend = (((reloc_value - sym->st_value)
                           + sym->st_value << 8)
                           | 0xe9);
        printf("[+] Rela entry for symbol: %s was modified\n", symbol_name);
        return true;
    }
}
```

Detection

- Check for relocations being applied at offset 0 on PIE binaries.
- This technique requires the text segment to be writable

```
00000021ffc8 000700000006 R_X86_64_GLOB_DAT 0000000000000000 free@GLIBC_2.2.5 + 0
00000021ffd0 000b00000006 R_X86_64_GLOB_DAT 0000000000000000 _ITM_deregisterTMClone + 0
00000021ffd8 003800000006 R_X86_64_GLOB_DAT 0000000000000000 __libc_start_main@GLIBC_2.2.5 + 0
000000000000 004400000001 R_X86_64_64 000000000001e240 __gmon_start__ + 5850e9
00000021ffe8 004e00000006 R_X86_64_GLOB_DAT 0000000000000000 malloc@GLIBC_2.2.5 + 0
00000021fff0 006e00000006 R_X86_64_GLOB_DAT 0000000000000000 _ITM_registerTMCloneTa + 0
00000021fff8 007200000006 R_X86_64_GLOB_DAT 0000000000000000 __cxa_finalize@GLIBC_2.2.5 + 0
000000220280 007700000005 R_X86_64_COPY 0000000000220280 __progname@GLIBC_2.2.5 + 0
000000220288 009300000005 R_X86_64_COPY 0000000000220288 stdout@GLIBC_2.2.5 + 0
```

Hiding Constructors in Dynamic and Static executables

Let's imagine we are an attacker and we would like to execute a constructor before reaching our main payload.

Constructors are easy to spot if we analyse the dynamic segment in dynamic executables

```
Dynamic section at offset 0xdc0 contains 27 entries:
```

Tag	Type	Name/Value
0x0000000000000001	(NEEDED)	Shared library: [libc.so.6]
0x000000000000000c	(INIT)	0x528
0x000000000000000d	(FINI)	0x7d4
0x0000000000000019	(INIT_ARRAY)	0x200db0
0x000000000000001b	(INIT_ARRAYSZ)	8 (bytes)
0x000000000000001a	(FINI_ARRAY)	0x200db8
0x000000000000001c	(FINI_ARRAYSZ)	8 (bytes)

Hiding Constructors in Dynamic and Static executables

- Would there be a way to hide used constructors from plain sight?
- Yes. We can we relocations

Technique Overview



Abusing IRELATIV relocations on static ELF binaries

- One could guess that hijacking Statically linked binaries would be as trivial as for dynamically linked binaries.
- However, this is not the case since statically linked executables only support one type of relocations, those being R_IRELATIVE.

What are IRELATIV relocations

From <https://sites.google.com/site/x32abi/documents/ifunc.txt>

This relocation is similar to `R_*_RELATIVE` except that the value used in this relocation is the program address returned by the function, which takes no arguments, at the address of the result of the corresponding `R_*_RELATIVE` relocation as specified in the processor-specific ABI.

The purpose of this relocation to avoid name lookup for locally defined `STT_GNU_IFUNC` symbols at load-time.


```
#define ELF_MACHINE_IREL 1
```

/glibc/sysdeps/x86_64/dl-irel.h

```
static inline ElfW(Addr)
__attribute__((always_inline))
elf_ifunc_invoke (ElfW(Addr) addr)
{
    return ((ElfW(Addr) (*) (void)) (addr)) ();
}

static inline void
__attribute__((always_inline))
elf_irela (const ElfW(Rela) *reloc)
{
    ElfW(Addr) *const reloc_addr = (void *) reloc->r_offset;
    const unsigned long int r_type = ELFW(R_TYPE) (reloc->r_info);

    if (__glibc_likely (r_type == R_X86_64_IRELATIVE))
    {
        ElfW(Addr) value = elf_ifunc_invoke(reloc->r_addend);
        *reloc_addr = value;
    }
    else
        __libc_fatal ("Unexpected reloc type in static binary.\n");
}
```

```

#define ELF_MACHINE_IRELAT 1

static inline ElfW(Addr)
__attribute__((always_inline))
elf_ifunc_invoke (ElfW(Addr) addr)
{
    return ((ElfW(Addr) (*)(void)) (addr)) ();
}

static inline void
__attribute__((always_inline))
elf_irela (const ElfW(Rela) *reloc)
{
    ElfW(Addr) *const reloc_addr = (void *) reloc->r_offset;
    const unsigned long int r_type = ELFW(R_TYPE) (reloc->r_info);

    if (__glibc_likely (r_type == R_X86_64_IRELATIVE))
    {
        ElfW(Addr) value = elf_ifunc_invoke(reloc->r_addend);
        *reloc_addr = value;
    }
    else
        __libc_fatal ("Unexpected reloc type in static binary.\n");
}

```

```
#define ELF_MACHINE_IRELA      1

static inline ElfW(Addr)
__attribute__((always_inline))
elf_ifunc_invoke (ElfW(Addr) addr)
{
    return ((ElfW(Addr) (*) (void)) (addr)) ();
}

static inline void
__attribute__((always_inline))
elf_irela (const ElfW(Rela) *reloc)
{
    ElfW(Addr) *const reloc_addr = (void *) reloc->r_offset;
    const unsigned long int r_type = ELFW(R_TYPE) (reloc->r_info);

    if (__glibc_likely (r_type == R_X86_64_IRELATIVE))
    {
        ElfW(Addr) value = elf_ifunc_invoke(reloc->r_addend);
        *reloc_addr = value;
    }
    else
        __libc_fatal ("Unexpected reloc type in static binary.\n");
}
```

How these ifuncs look like

```
[0x00423480]> ir~0x00423480
vaddr=0x006b9018 paddr=0x000b9018 type=SET_64 0x00423480 (ifunc)
[0x00423480]> pdf
      ;-- strcpy:
/ (fcn) sym.strcpy_ifunc 43
|   sym.strcpy_ifunc ();
|       0x00423480      f605a5892900.  test byte [0x006bbe2c], 0x10
|       0x00423487      488d05020602.  lea rax, sym.__strcpy_sse2_unaligned
|       ,=< 0x0042348e      7519      jne 0x4234a9
|       | 0x00423490      f6055a892900.  test byte [0x006bbdf1], 2
|       | 0x00423497      488d0572d401.  lea rax, sym.__strcpy_sse2
|       | 0x0042349e      488d152bd601.  lea rdx, sym.__strcpy_ssse3
|       | 0x004234a5      480f45c2      cmovne rax, rdx
|       `-> 0x004234a9      f3c3      ret
[0x00423480]> □
```

An approach for IRELATIV Reloc hijacking

- 1 - Emulate all/specific ifuncs and populate it correspondent GOT entry with the function result.
- 2 - All IRELATIV relocations now are free to be tampered.
- 3 - Create a fake ifunc function that returns the address of some arbitrary malicious code.
- 4 - change an IRELATIV relocation entry so that the addend points to the address of the fake ifunc function.
- 5 - change the offset of the relocation to point to a known constructor/destructor data structure.
- 6 - ???
- 7 - profit.

Implementation with Radare2

```
1 import r2pipe
2 import sys
3
4 def main(file_path):
5     r = r2pipe.open(file_path, flags=['-w'])
6     relocs_value = r.cmd('ir~[6]').split()
7     relocs_offset = r.cmd('ir~[1]').split()
8
9     print("[+] Resolving all ifunc relocations");
10    for value, offset in zip(relocs_value, relocs_offset):
11        r.cmd("s %s" % value)
12        r.cmd("aei")
13        r.cmd("aeip")
14        r.cmd("aesuo ret")
15        resolved_ifunc = r.cmd("aer rax").strip()
16        print("[+] Writing %s at %s" % (resolved_ifunc, offset))
17        r.cmd("wv8 %s @%s" % (resolved_ifunc, offset))
18        print r.cmd("pxq 8 @ %s" % offset);
19        r.cmd("aei-")
20
21    dummy_pointer = r.cmd("is~__fini_array_start~[2]").strip()
22    payload_address = r.cmd("is~fake_ireloc~[2]").strip()
23
24    print("[+] Creating new IRELATIV relocation (%s, %s)" % (dummy_pointer, payload_address))
25    r.cmd("s section..rela.plt")
26    r.cmd('wv8 %s @ $$ ' % dummy_pointer);
27    r.cmd('wv8 %s @ $$+0x10' % payload_address);
28
29 if __name__ == '__main__':
30     main(sys.argv[1])
```

Demo

Detection on dynamic executables

Before

Relocation section '.rela.dyn' at offset 0x438 contains 8 entries:

Offset	Info	Type	Sym. Value	Sym. Name + Addend
000000200db0	000000000008	R_X86_64_RELATIVE		680
000000200db8	000000000008	R_X86_64_RELATIVE		640
000000201008	000000000008	R_X86_64_RELATIVE		201008
000000200fd8	000100000006	R_X86_64_GLOB_DAT	0000000000000000	_ITM_deregisterTMClone + 0
000000200fe0	000300000006	R_X86_64_GLOB_DAT	0000000000000000	__libc_start_main@GLIBC_2.2.5 + 0
000000200fe8	000400000006	R_X86_64_GLOB_DAT	0000000000000000	__gmon_start__ + 0
000000200ff0	000600000006	R_X86_64_GLOB_DAT	0000000000000000	_ITM_registerTMCloneTa + 0
000000200ff8	000700000006	R_X86_64_GLOB_DAT	0000000000000000	__cxa_finalize@GLIBC_2.2.5 + 0

Relocation section '.rela.plt' at offset 0x4f8 contains 2 entries:

Offset	Info	Type	Sym. Value	Sym. Name + Addend
000000200fc8	000200000007	R_X86_64_JUMP_SLO	0000000000000000	puts@GLIBC_2.2.5 + 0
000000200fd0	000500000007	R_X86_64_JUMP_SLO	0000000000000000	mprotect@GLIBC_2.2.5 + 0

Detection on dynamic executables

After

Relocation section '.rela.dyn' at offset 0x438 contains 8 entries:

Offset	Info	Type	Sym. Value	Sym. Name + Addend
000000200db0	000000000008	R_X86_64_RELATIVE		6ac
000000200db8	000000000008	R_X86_64_RELATIVE		640
000000201008	000000000008	R_X86_64_RELATIVE		201008
000000200fd8	000100000006	R_X86_64_GLOB_DAT	0000000000000000	_ITM_deregisterTMClone + 0
000000200fe0	000300000006	R_X86_64_GLOB_DAT	0000000000000000	__libc_start_main@GLIBC_2.2.5 + 0
000000200fe8	000400000006	R_X86_64_GLOB_DAT	0000000000000000	__gmon_start__ + 0
000000200ff0	000600000006	R_X86_64_GLOB_DAT	0000000000000000	_ITM_registerTMCloneTa + 0
000000200ff8	000700000006	R_X86_64_GLOB_DAT	0000000000000000	__cxa_finalize@GLIBC_2.2.5 + 0

Relocation section '.rela.plt' at offset 0x4f8 contains 2 entries:

Offset	Info	Type	Sym. Value	Sym. Name + Addend
000000200fc8	000200000007	R_X86_64_JUMP_SLO	0000000000000000	puts@GLIBC_2.2.5 + 0
000000200fd0	000500000007	R_X86_64_JUMP_SLO	0000000000000000	mprotect@GLIBC_2.2.5 + 0

Detection on dynamic executables

After

Relocation section '.rela.dyn' at offset 0x438 contains 8 entries:

Offset	Info	Type	Sym. Value	Sym. Name + Addend
000000200db0	000000000008	R_X86_64_RELATIVE		6ac
000000200db8	000000000008	R_X86_64_RELATIVE		640
000000201008	000000000008	R_X86_64_RELATIVE		201008
000000200fd8	000100000006	R_X86_64_GLOB_DAT	0000000000000000	_ITM_deregisterTMClone + 0
000000200fe0	000300000006	R_X86_64_GLOB_DAT	0000000000000000	__libc_start_main@GLIBC_2.2.5 + 0
000000200fe8	000400000006	R_X86_64_GLOB_DAT	0000000000000000	__gmon_start__ + 0
000000200ff0	000600000006	R_X86_64_GLOB_DAT	0000000000000000	_ITM_registerTMCloneTa + 0
000000200ff8	000700000006	R_X86_64_GLOB_DAT	0000000000000000	__cxa_finalize@GLIBC_2.2.5 + 0

Relocation section '.rela.plt' at offset 0x4f8 contains 2 entries:

Offset	Info	Type	Sym. Value	Sym. Name + Addend
000000200fc8	000200000007	R_X86_64_JUMP_SLO	0000000000000000	puts@GLIBC_2.2.5 + 0
000000200fd0	000500000007	R_X86_64_JUMP_SLO	0000000000000000	mprotect@GLIBC_2.2.5 + 0

Detection on static executables

- Check where IRELATIV relocations points to.
- They all should point to STT_IFUNC type symbols.

Problem: these can be spotted thanks to the symbol table (STT_IFUNC symbols) however string/symbol tables can completely removed/instrumented in statically linked executables.

Experimental Practices



Removing .dynstr string table

- .Dynstr string table is the string table of .dynsym symbol table
- These two artifacts are needed on runtime by the Dynamic Linker
- Usually other string tables can be stripped but not .dynstr
- All dynamically linked executables have them

The following technique will explain an approach on how this string table can be removed to leverage an anti-analysis technique

How imports get resolved

- Imports get resolved thanks to the PLT mechanism implemented in most dynamic executable ELF files.
- There is a way for a dynamic executable to not contain a PLT, although is not common.

How imports get resolved

```

;-- section..plt:
0x00003770    ff35cac42100    push qword [0x0021fc40]
0x00003776    ff25ccc42100    jmp qword [0x0021fc48]
0x0000377c    0f1f4000        nop dword [rax]
;-- imp.__ctype_toupper_loc:
0x00003780    ff25cac42100    jmp qword reloc.__ctype_toupper_loc
0x00003786    6800000000      push 0
0x0000378b    e9e0ffffff      jmp section..plt
;-- imp.__uflow:
0x00003790    ff25c2c42100    jmp qword reloc.__uflow
0x00003796    6801000000      push 1
0x0000379b    e9d0ffffff      jmp section..plt
;-- imp.getenv:
0x000037a0    ff25bac42100    jmp qword reloc.getenv
0x000037a6    6802000000      push 2
0x000037ab    e9c0ffffff      jmp section..plt
;-- imp.sigprocmask:
0x000037b0    ff25b2c42100    jmp qword reloc.sigprocmask
0x000037b6    6803000000      push 3
0x000037bb    e9b0ffffff      jmp section..plt
[0x00003770]> 
```


How imports get resolved

First 3 GOT entries have reserved values.

- GOT[0] - Virtual address of the image's DYNAMIC segment
- GOT[1] - Virtual address of link_map
- GOT[2] - Virtual address of _dl_runtime_resolve

```

;-- section..plt:
0x00003770    ff35cac42100    push qword [0x0021fc40] GOT[1] - link_map
0x00003776    ff25ccc42100    jmp qword [0x0021fc48]  GOT[2] - _dl_runtime_resolve
0x0000377c    0f1f4000       nop dword [rax]

;-- imp.__ctype_toupper_loc:
0x00003780    ff25cac42100    jmp qword reloc.__ctype_toupper_loc
0x00003786    6800000000     push 0
0x0000378b    e9e0ffffff     jmp section..plt

```

Relocation index in .rela.plt

How imports get resolved

```
result = _dl_lookup_symbol_x (strtab + sym->st_name, l, &sym, l->l_scope,
                             version, ELF_RTYPE_CLASS_PLT, flags, NULL);

/* We are done with the global scope. */
if (!RTLD_SINGLE_THREAD_P)
    THREAD_GSCOPE_RESET_FLAG ();

#ifdef RTLD_FINALIZE_FOREIGN_CALL
    RTLD_FINALIZE_FOREIGN_CALL;
#endif

/* Currently result contains the base load address (or link map)
   of the object that defines sym. Now add in the symbol
   offset. */
value = DL_FIXUP_MAKE_VALUE (result,
                             SYMBOL_ADDRESS (result, sym, false));
```

How imports get resolved

```
result = _dl_lookup_symbol_x (strtab + sym->st_name, l, &sym, l->l_scope,
                             version, ELF_RTYPE_CLASS_PLT, flags, NULL);

/* We are done with the global scope. */
if (!RTLD_SINGLE_THREAD_P)
    THREAD_GSCOPE_RESET_FLAG ();

#ifdef RTLD_FINALIZE_FOREIGN_CALL
    RTLD_FINALIZE_FOREIGN_CALL;
#endif

/* Currently result contains the base load address (or link map)
   of the object that defines sym. Now add in the symbol
   offset. */
value = DL_FIXUP_MAKE_VALUE (result,
                             SYMBOL_ADDRESS (result, sym, false));
```

The Challenge

- Is there a way to resolve symbols without the need to rely on their names? Yes
- How? Creating a DT_HASH/DT_GNU_HASH resolver

Dynamic section at offset 0x27e68 contains 19 entries:

Tag	Type	Name/Value
0x000000000000000e	(SONAME)	Library soname: [ld-linux-x86-64.so.2]
0x0000000000000004	(HASH)	0x1f0
0x0000000006ffffef5	(GNU_HASH)	0x2c8
0x0000000000000005	(STRTAB)	0x6f0
0x0000000000000006	(SYMTAB)	0x3c0

The approach (experimental WIP)

1. Hash all dynamic symbol strings for subject binary and remove .dynstr.
2. Save all hashes in a hash-table.
3. Make sure the hashes are placed in the hash table in the same order as the relocation indexes per dynamic symbol.
4. Inject a custom hash resolver into target binary along with the hashing table.
5. Inject constructor that replaces the value of GOT[2] for our resolver on runtime.
6. Profit.

Demo

Outcome

Symbol table '.dynsym' contains 546 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	0000000000000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
2:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
3:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
4:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
5:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
6:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
7:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
8:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
9:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
10:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
11:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
12:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
13:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
14:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
15:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
16:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	
17:	0000000000000000	0	FUNC	GLOBAL	DEFAULT	UND	

Outcome

```
got: 0000000000280B78 nullsub_15_ptr dq offset nullsub_15
got: 0000000000280B80 nullsub_16_ptr dq offset nullsub_16
got: 0000000000280B88 nullsub_17_ptr dq offset nullsub_17
got: 0000000000280B90 nullsub_18_ptr dq offset nullsub_18
got: 0000000000280B98 nullsub_19_ptr dq offset nullsub_19
got: 0000000000280BA0 nullsub_20_ptr dq offset nullsub_20
got: 0000000000280BA8 nullsub_21_ptr dq offset nullsub_21
got: 0000000000280BB0 nullsub_22_ptr dq offset nullsub_22
got: 0000000000280BB8 nullsub_23_ptr dq offset nullsub_23
got: 0000000000280BC0 nullsub_24_ptr dq offset nullsub_24
got: 0000000000280BC8 nullsub_25_ptr dq offset nullsub_25
got: 0000000000280BD0 nullsub_26_ptr dq offset nullsub_26
got: 0000000000280BD8 nullsub_27_ptr dq offset nullsub_27
got: 0000000000280BE0 nullsub_28_ptr dq offset nullsub_28
got: 0000000000280BE8 nullsub_29_ptr dq offset nullsub_29
got: 0000000000280BF0 nullsub_30_ptr dq offset nullsub_30
got: 0000000000280BF8 nullsub_31_ptr dq offset nullsub_31
got: 0000000000280C00 nullsub_32_ptr dq offset nullsub_32
got: 0000000000280C08 nullsub_33_ptr dq offset nullsub_33
got: 0000000000280C10 nullsub_34_ptr dq offset nullsub_34
got: 0000000000280C18 nullsub_35_ptr dq offset nullsub_35
got: 0000000000280C20 nullsub_36_ptr dq offset nullsub_36
```


Outcome

got: 0000000000280B78 nullsub_15_ptr dq offset nullsub_15
got: 0000000000280B80 nullsub_16_ptr dq offset nullsub_16
got: 0000000000280B88 nullsub_17_ptr dq offset nullsub_17
got: 0000000000280B90 nullsub_18_ptr dq offset nullsub_18
got: 0000000000280C00 nullsub_29_ptr dq offset nullsub_29
got: 0000000000280C08 nullsub_33_ptr dq offset nullsub_33
got: 0000000000280C10 nullsub_34_ptr dq offset nullsub_34
got: 0000000000280C18 nullsub_35_ptr dq offset nullsub_35
got: 0000000000280C20 nullsub_36_ptr dq offset nullsub_36

IDA View-AHex View-1StructuresEnumsImports

Address	Ordinal	Name	Library
0000000000280C00		nullsub_29_ptr	
0000000000280C08		nullsub_33_ptr	
0000000000280C10		nullsub_34_ptr	
0000000000280C18		nullsub_35_ptr	
0000000000280C20		nullsub_36_ptr	

Outcome

```
[0x0000d090]> ii
```

```
[Imports]
```

Num	Vaddr	Bind	Type	Name	0x0000d09f	4883ec28	sub rsp, 0x28
					0x0000d0a3	4c8b26	mov r12, qword [rsi]
1	0x00000000	GLOBAL	FUNC		0x0000d0a6	be2f000000	mov esi, 0x2f
2	0x00000000	GLOBAL	FUNC		0x0000d0ab	64488b042528.	mov rax, qword fs:[0x28]
3	0x00000000	GLOBAL	FUNC		0x0000d0b4	4889442418	mov qword [rsp + 0x18], rax
4	0x00000000	GLOBAL	FUNC		0x0000d0b9	31c0	xor eax, eax
5	0x00000000	GLOBAL	FUNC		0x0000d0bb	4c89e7	mov rdi, r12
6	0x00000000	GLOBAL	FUNC		0x0000d0be	e85ddbffff	call 0xac20
7	0x00000000	GLOBAL	FUNC		0x0000d0c3	488d5001	lea rdx, [rax + 1]
8	0x00000000	GLOBAL	FUNC		0x0000d0c7	4885c0	test rax, rax
9	0x00000000	GLOBAL	FUNC		0x0000d0ca	4c0f45e2	cmovne r12, rdx
10	0x00000000	GLOBAL	FUNC		0x0000d0ce	4531ff	xor r15d, r15d
11	0x00000000	GLOBAL	FUNC		0x0000d0d1	83fb01	cmp ebx, 1
12	0x00000000	GLOBAL	FUNC				
13	0x00000000	GLOBAL	FUNC				
14	0x00000000	GLOBAL	FUNC				
15	0x00000000	GLOBAL	FUNC				
16	0x00000000	GLOBAL	FUNC				
17	0x00000000	GLOBAL	FUNC				

Conclusion

- The ELF file format is very flexible.
- Parsing anomalies are easy.
- Code injection may start becoming more common as more complex implants arise.
- Relocations can be used to hide conventional control flow hijacking methodologies.
- ELF files can be easily abused for anti-analysis purposes in ways today are not being used.
- Expecting an increase in Linux malware complexity in coming years.

Questions?