Object Diversification with the help of R2

Alex Gaines, R2con 2019

This project is sponsored by the Intelligence Advanced Research Projects Activity (IARPA) under contract N66001-17-C-2003

Company Overview

- Founded 2009 by Jason Syversen
- 30+ Employees
- Most possess advanced degrees
- All have current (or in process for) clearances

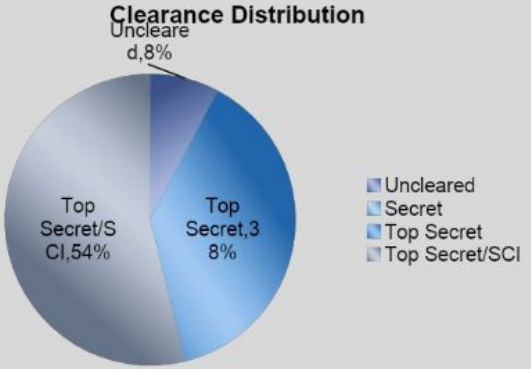

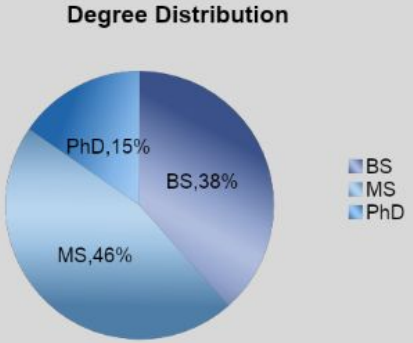

About me:

Alex Gaines

Alex.Gaines@siegetechnologies.com

Siege Technologies.

I like bins.

A few R2 commits.

Also check out WPICTF.xyz

Automated Software Diversification - types

- Load Time
  - ASLR
    - On some embedded systems, may not be enabled or possible.
- Compile Time
  - Usually operate on AST or Source
  - Slow - must compile
- **Binary**
  - Can’t be very advanced.
  - Tends to break under certain cases.

Software Diversification. - Why

- Defense
  - Increase attacker workload.
    - Protect against Return Oriented Programming and other techniques.
  - Obfuscate.
  - Anti Piracy
- Offense
  - Avoid signatures.
    - AV/Intrusion detection avoidance.
  - Obfuscate

Goal

- Protect Kernel Modules from ROP attacks.
- Diversification
  - Every boot, Kernel Modules are different
  - Change ROP offsets, stack offsets/layout

Limitations

- Fast
- Small
- Reliable
- Operate on compiled kernel modules, not source

Quick Elf Overview

Executable and Linkable Format

Used often in Linux executables, objects, libraries, etc.

Consists of headers and sections.

We need to worry about Symbols and Relocations

R2 has elf code

But can also use other libraries, make your own, etc.

We made our own.

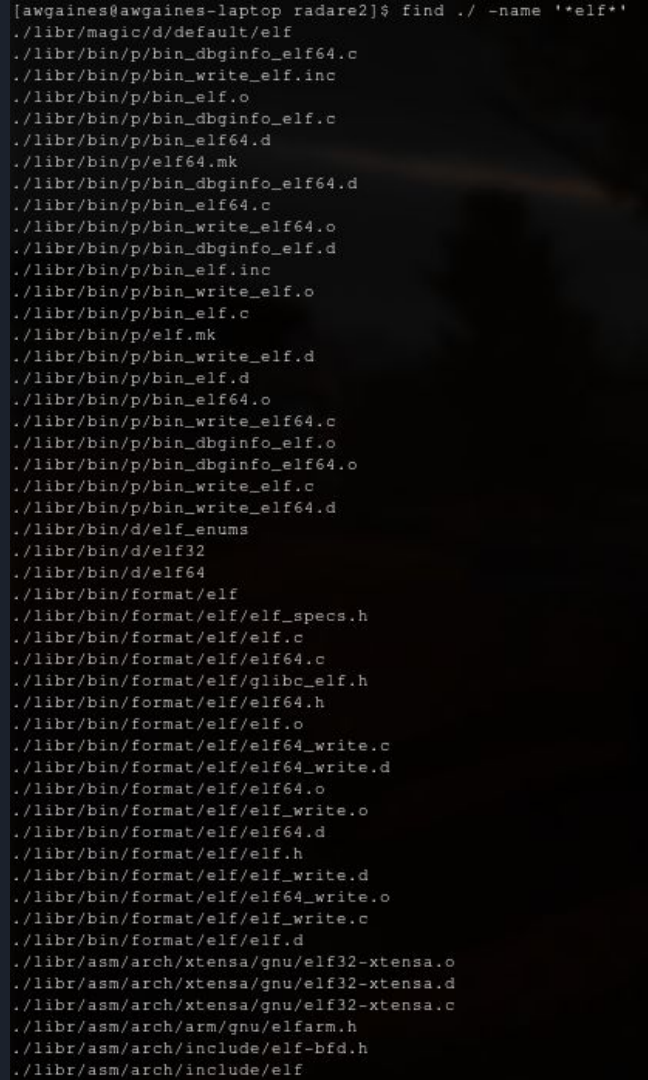

Easy Start

- Shuffle elf sections
  - -ffunction-sections splits every function into a section
  - Just update some ELF metadata indices.
  - OpenBSD does something similar
    - Re-Links LibC's from .a files in a random order at every boot.
  - Limited # of permutations
    - More sections = better results

Easy Start

- Junk data at the end of sections
  - Random # of HLT instructions at the end of .text* sections
  - Comes after all the code
  - Changes offsets between sections.
    - But relative offsets inside sections are unchanged.

Why these aren’t enough

- Only change offsets between sections, not inside sections.
- Limited permutations.
- Does not protect stack in any way.

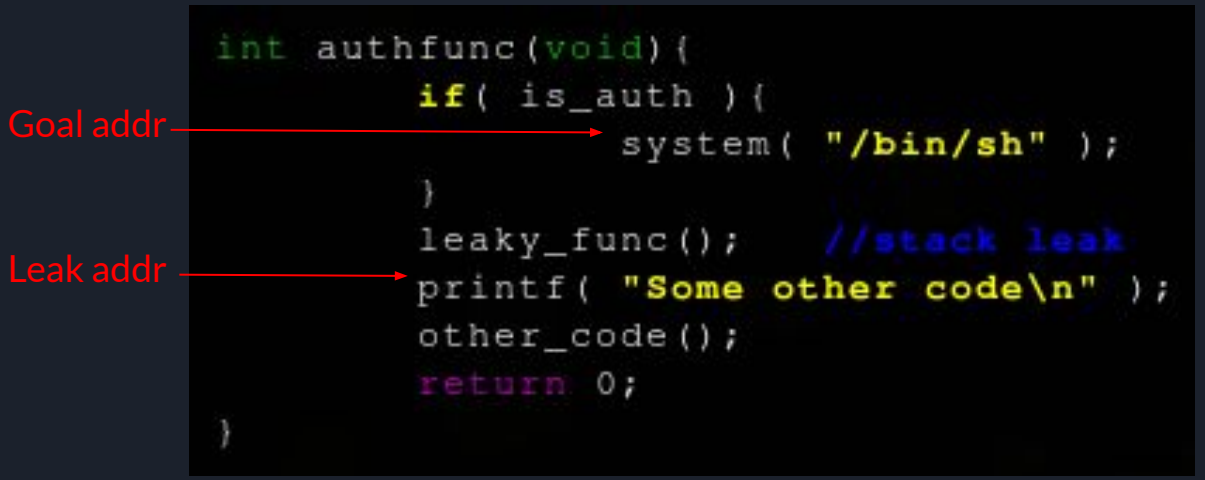

Let’s mess code up.

We need some help

- Change offsets between Leaked pointer and useful gadget.
  - Code insertion.
  - Can’t use “Trampoline style” dynamic code insertion. - SKORPIO.
- Want to modify the code itself.
  - Actually insert code, not just modify existing.
  - We need to “analyze” our sections.

Why Radare2?

- Fast
- Low dependencies (No python required!)
- Compiles core to LibR
- Has some neat features (Esil)
- **I know how to use it**

Why libr and not r2pipe?

- **Fast**, a lot less overhead.
- Easier to use a real library than to interface with a text api.
  - At least for C.

Analysis overview.

Disassemble every executable section using r_anal_op() loops

```
R_API int r_anal_op(RAnal *anal, RAnalOp *op, ut64 addr,
        const ut8 *data, int len, RAnalOpMask mask);
```

What this gives us:

- For each function (Symbol):
  - An ordered list of RAnalOps & their locations.
    - Instruction Type.
    - Jump targets.
    - Instruction opcode vs data size.
    - Registers used, memory access, etc.
    - **ESIL too!**

RAnalOp Is Super Useful.

```
typedef struct r_anal_op_t {
        char *mnemonic; /* mnemonic.. it actually contains the args too. we should replace rasm with this */
        ut64 addr;      /* address */
        ut32 type;      /* type of opcode */
        ut64 prefix;    /* type of opcode prefix (rep,lock,...) */
        ut32 type2;     /* used by java */
        int group;      /* is fpu, is privileged, mmx, etc */
        int stackop;    /* operation on stack? */
        int cond;       /* condition type */
        int size;       /* size in bytes of opcode */
        int nopcode;    /* number of bytes representing the opcode (not the arguments) TODO: find better name */
        int cycles;     /* cpu-cycles taken by instruction */
        int failcycles; /* conditional cpu-cycles */
        int family;     /* family of opcode */
        int id;         /* instruction id */
        bool eob;       /* end of block (boolean) */
        bool sign;      /* operates on signed values, false by default */
        /* Run N instructions before executing the current one */
        int delay;      /* delay N slots (mips, ..)*/
        ut64 jump;      /* true jmp */
        ut64 fail;      /* false jmp */
        int direction;  /* 1 = read, 2 = write, 4 = exec, 8 = reference,  */
        st64 ptr;       /* reference to memory */ /* XXX signed? */
        ut64 val;       /* reference to value */ /* XXX signed? */
        int ptrsize;    /* f.ex: zero extends for 8, 16 or 32 bits only */
        st64 stackptr;  /* stack pointer */
        int refptr;     /* if (0) ptr = "reference" else ptr = "load memory of refptr bytes" */
        RAnalVar *var;  /* local var/arg used by this instruction */
        RAnalValue *src[3];
        RAnalValue *dst;
        struct r_anal_op_t *next; // TODO deprecate
        RStrBuf esil;
        RStrBuf opex;
        const char *reg; /* destination register */
        const char *ireg; /* register used for indirect memory computation*/
        int scale;
        ut64 disp;
        RAnalSwitchOp *switch_op;
        RAnalHint hint;
        RAnalDataType datatype;
} RAnalOp;
```

Esil is great too

- Great for easily gathering instruction information.
  - strstr() is your greatest friend here.
- Not terribly great at performance, but it's useful to get something working.
  - Does this instruction use the stack? strstr(esil, "rsp")

```
0,rax,rax,&,==,$z,zf,:=,$p,pf,:=,$s,sf,:=,0,cf,:=,0,of,:=
```

“Cruftables”

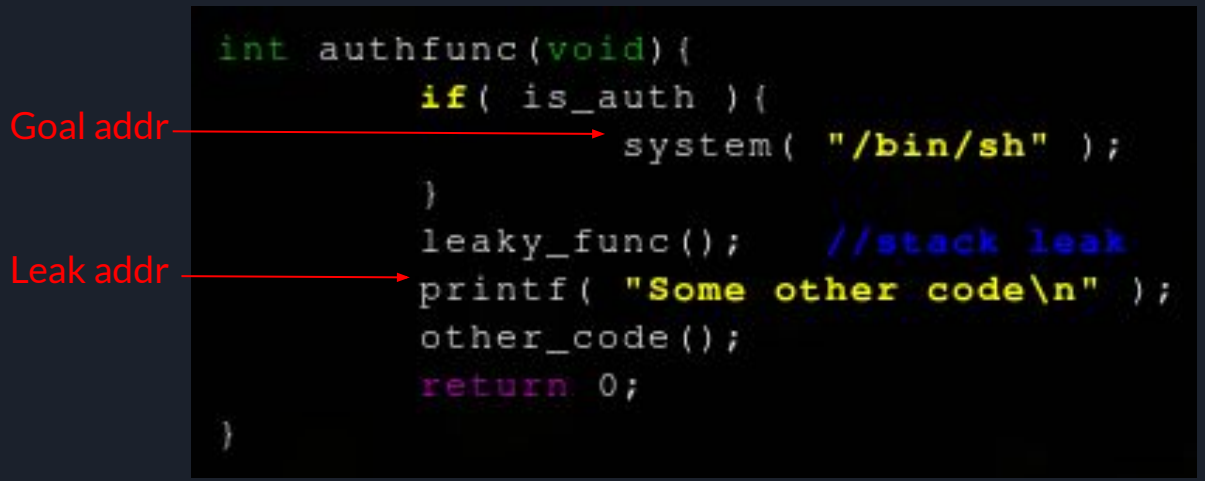
Goal addr
Leak addr
int authfunc(void){
if(is_auth){
system("/bin/sh");
}
leaky_func(); //stack leak
printf("Some other code\n");
other_code();
return 0;
}

“Cruftables”

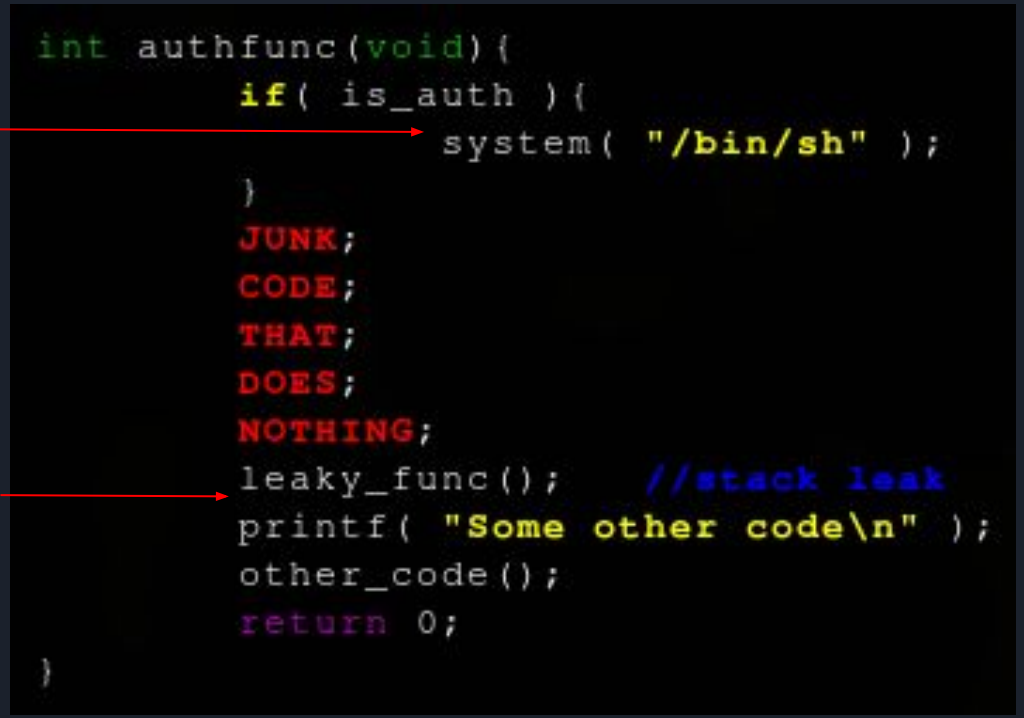
Goal addr
Leak addr
int authfunc(void){
if(is_auth){
system("/bin/sh");
}
JUNK;
CODE;
THAT;
DOES;
NOTHING;
leaky_func(); //stack leak
printf("Some other code\n");
other_code();
return 0;
}

What needs to be done?

Decide where to put a “cruftable”.

Decide what data is a “cruftable”?

Adjust symbols, relocs, etc to accommodate cruftable.

Where can we put a Cruftable?

- We have to insert in between instructions.
  - Can’t cut an instruction in half.
- Can’t put a cruftable in a “badzone”
  - Badzone consists of anywhere between a jump instruction and its jump target.
  - Not quite a basic block, but similar.

Badzone(s)

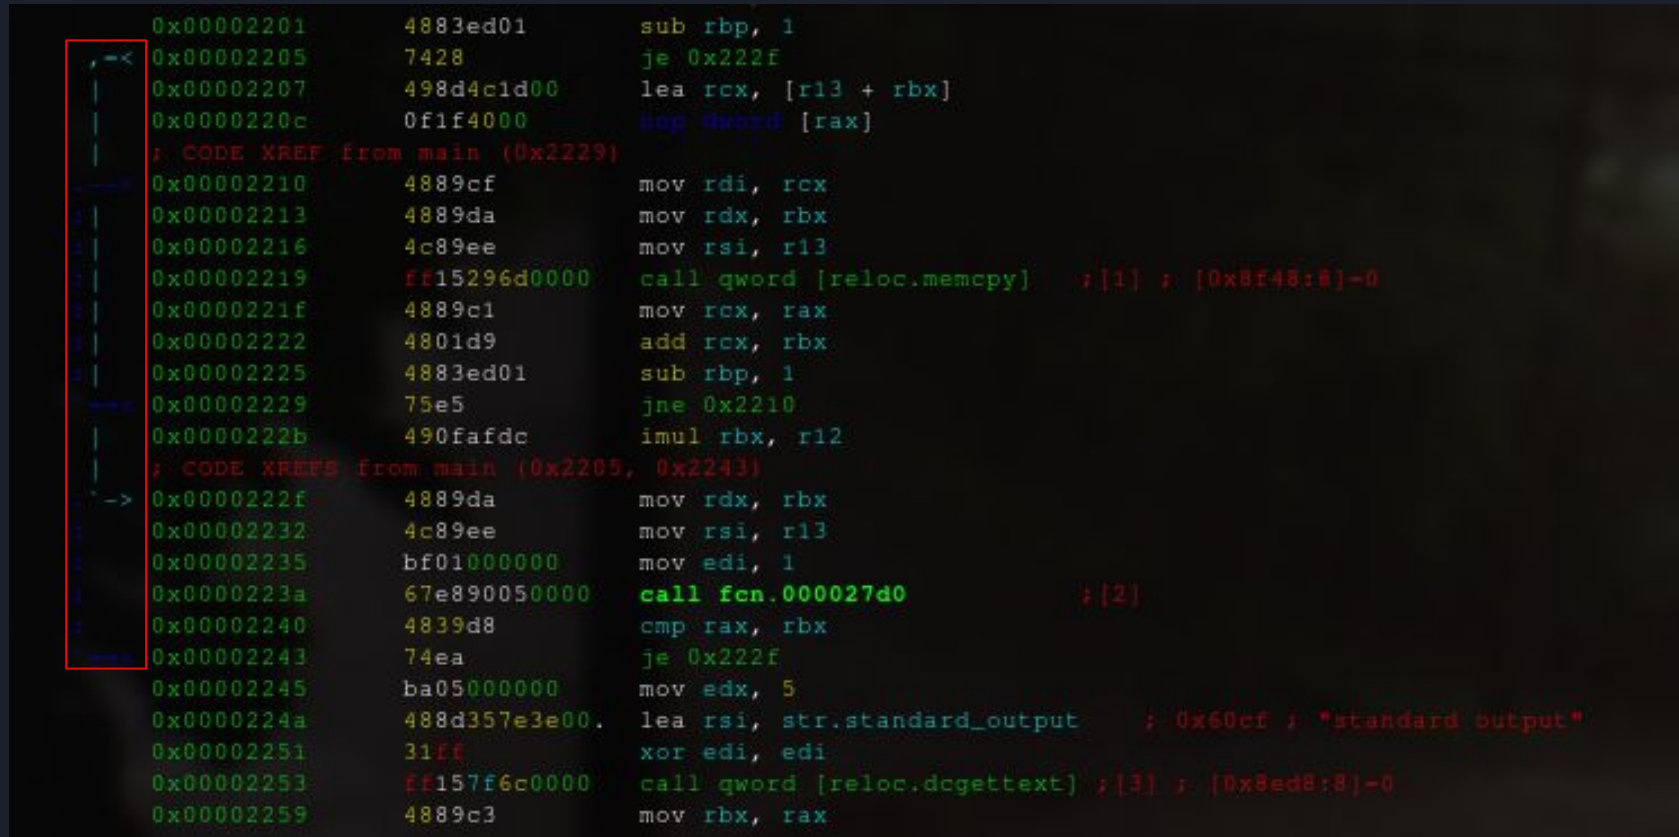
0x00002201 4883ed01 sub rbp, 1
0x00002205 7428 je 0x222f
0x00002207 498d4c1d00 lea rcx, [r13 + rbx]
0x0000220c 0f1f4000 nop dword [rax]
; CODE XREF from main (0x2229)
0x00002210 4889cf mov rdi, rcx
0x00002213 4889da mov rdx, rbx
0x00002216 4c89ee mov rsi, r13
0x00002219 ff15296d0000 call qword [reloc.memcpy] ;[1] ; [0x8f48:8]=0
0x0000221f 4889c1 mov rcx, rax
0x00002222 4801d9 add rcx, rbx
0x00002225 4883ed01 sub rbp, 1
0x00002229 75e5 jne 0x2210
0x0000222b 490fafdc imul rbx, r12
; CODE XREFS from main (0x2205, 0x2243)
0x0000222f 4889da mov rdx, rbx
0x00002232 4c89ee mov rsi, r13
0x00002235 bf01000000 mov edi, 1
0x0000223a 67e890050000 call fcn.000027d0 ;[2]
0x00002240 4839d8 cmp rax, rbx
0x00002243 74ea je 0x222f
0x00002245 ba05000000 mov edx, 5
0x0000224a 488d357e3e00. lea rsi, str.standard_output ; 0x60cf ; "standard output"
0x00002251 31ff xor edi, edi
0x00002253 ff157f6c0000 call qword [reloc.dcgettext] ;[3] ; [0x8ed8:8]=0
0x00002259 4889c3 mov rbx, rax

RAnalOp to Badzone.

- RAnalOp has a lot of useful data to use.
- Many ways to determine badzone-causing Instruction.
  - **Most important (RAnalOp.jump != -1 || RAnalOp.fail != -1)**
- Badzone = Bounds (instr.loc, instr.loc + instr.size, RAnalOp.jump, RAnalOp.fail)

Data in a cruftable.

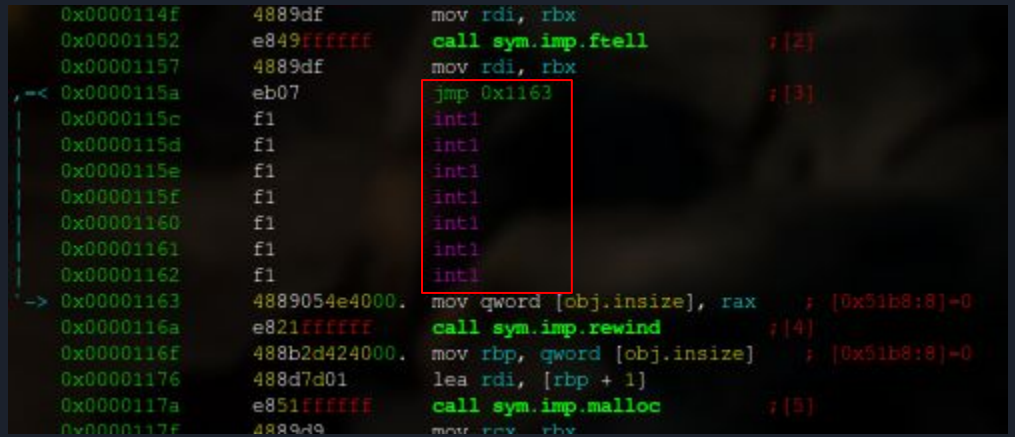

Data in a cruftable.

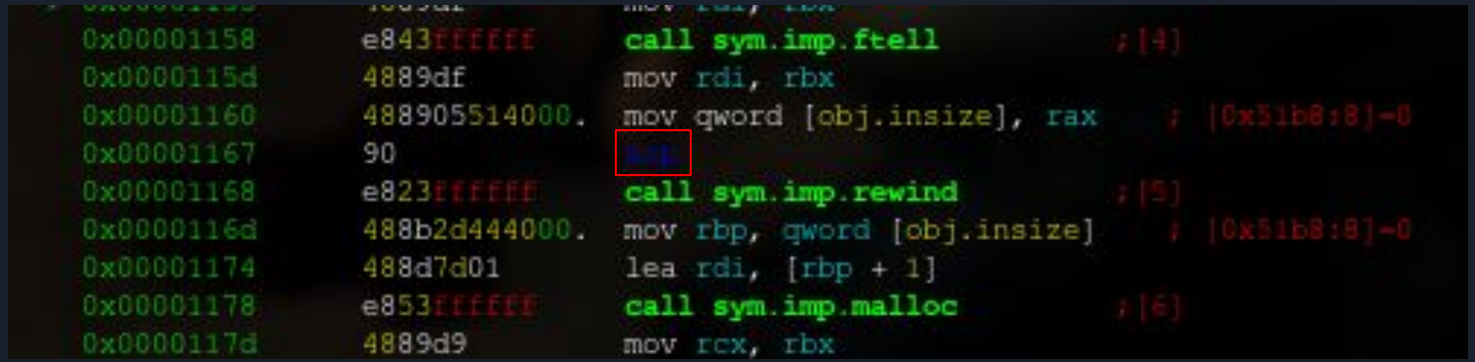

Cruftables can be long.

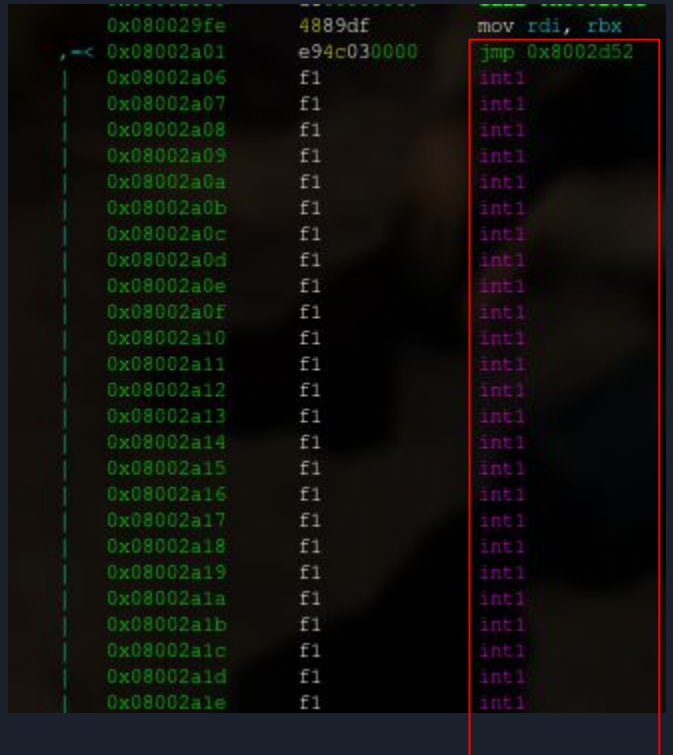

Elf data adjustment.

- Any Symbol that comes after a cruftable needs to be adjusted.
  - (For every Symbol: if Symbol is after cruftable, add cruftable size to Symbol)
- Any Reloc that comes after a cruftable needs to be adjusted.
  - (For every Reloc: if Reloc is after cruftable, add cruftable size to Reloc)

Elf data adjustment.

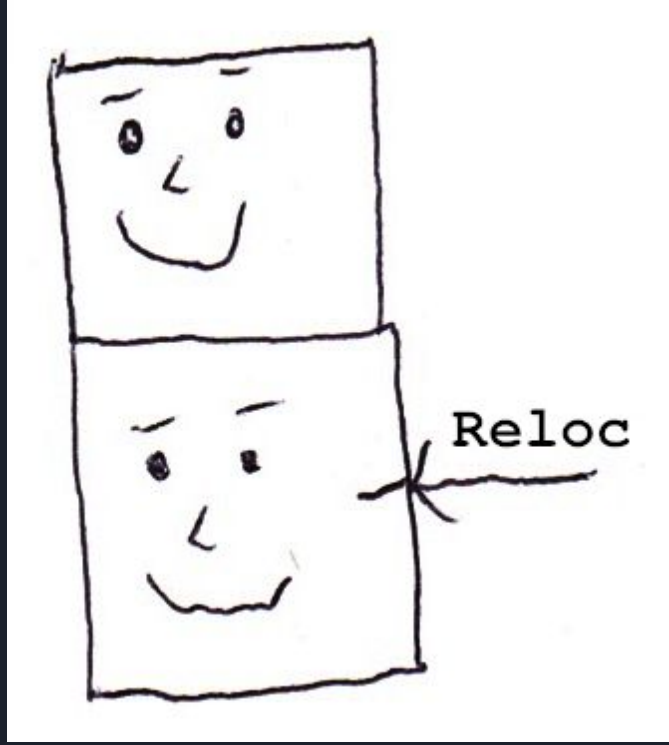

Elf data adjustment.

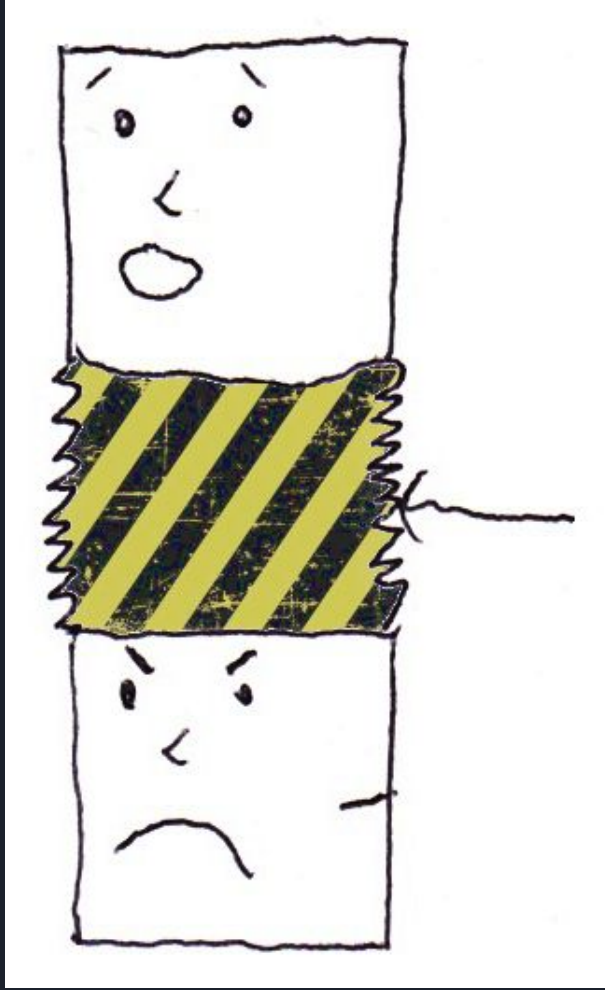

Elf data adjustment.

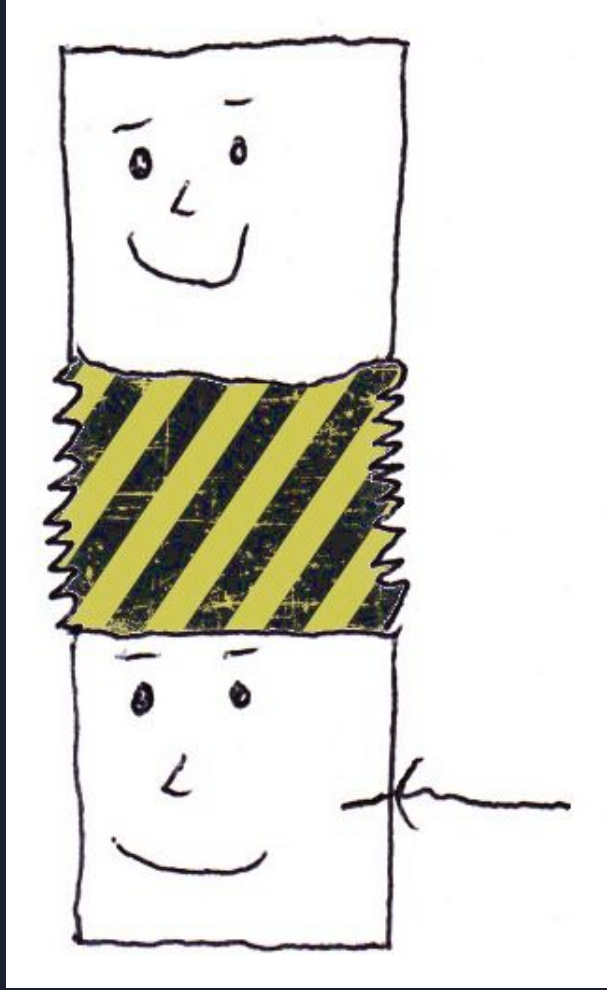

Split Cruftables

- When Inserting a cruftable: Randomly split into 2 cruftables.
  - Can be applied recursively.

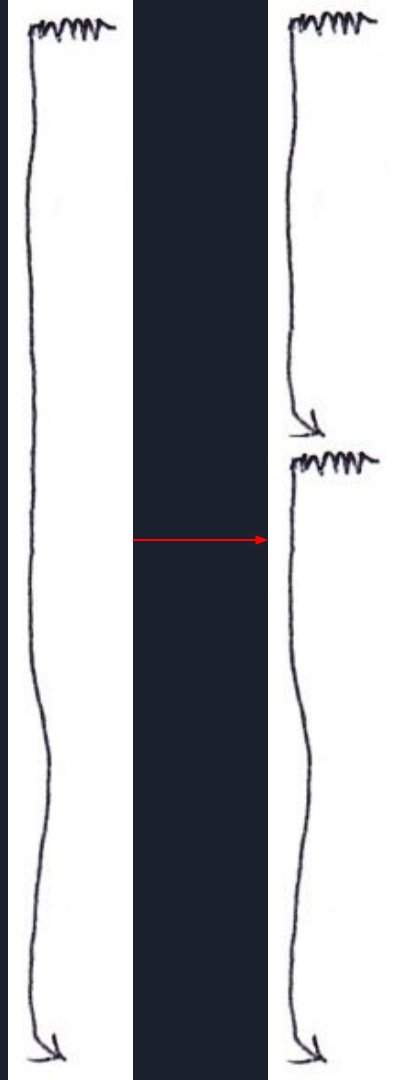

A split cruftable.

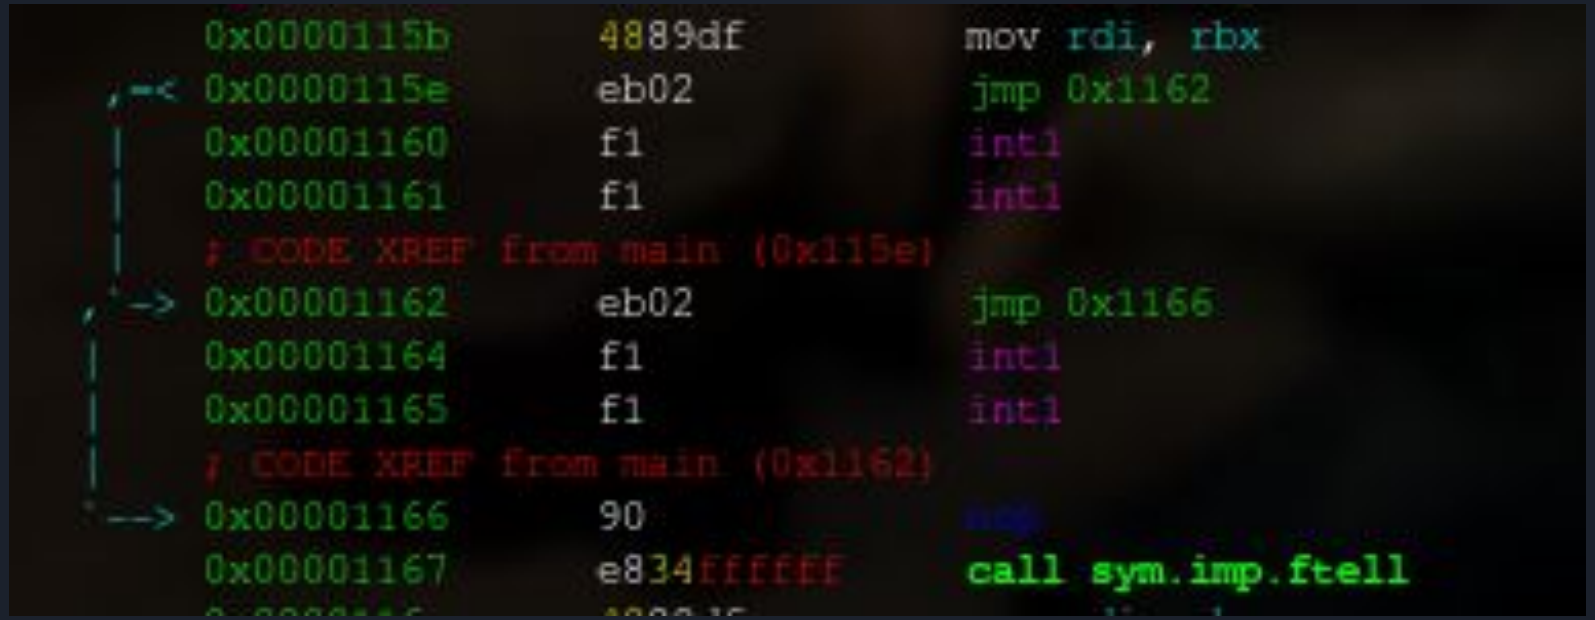

Confused Cruftables

- Instead of using jumps, use conditional jumps.
- “Fail” path must have a another cruftable inserted after it to make sure it correctly reaches the end.
  - Can be recursive, can be split.
  - Easily done recursively.
- When combined with split cruftables, it creates gnarly code.

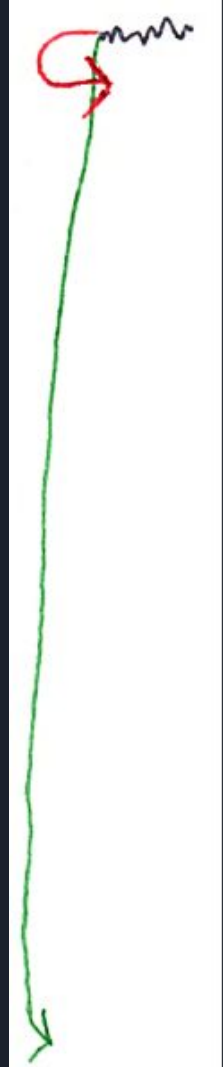

Confused Cruftables

- Instead of using jumps, use conditional jumps.
- “Fail” path must have a another cruftable inserted after it to make sure it correctly reaches the end.
  - Can be recursive, can be split.
  - Easily done recursively.
- When combined with split cruftables, it creates gnarly code.

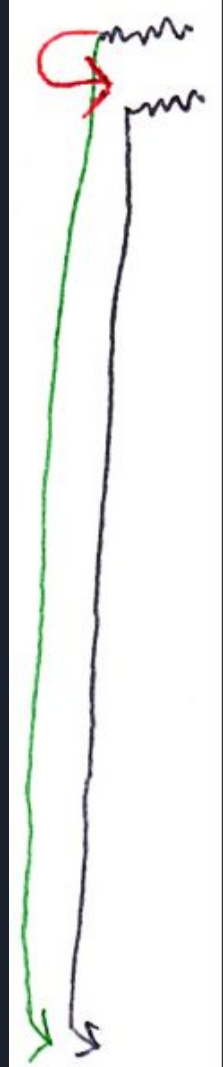

Confused and split cruftable

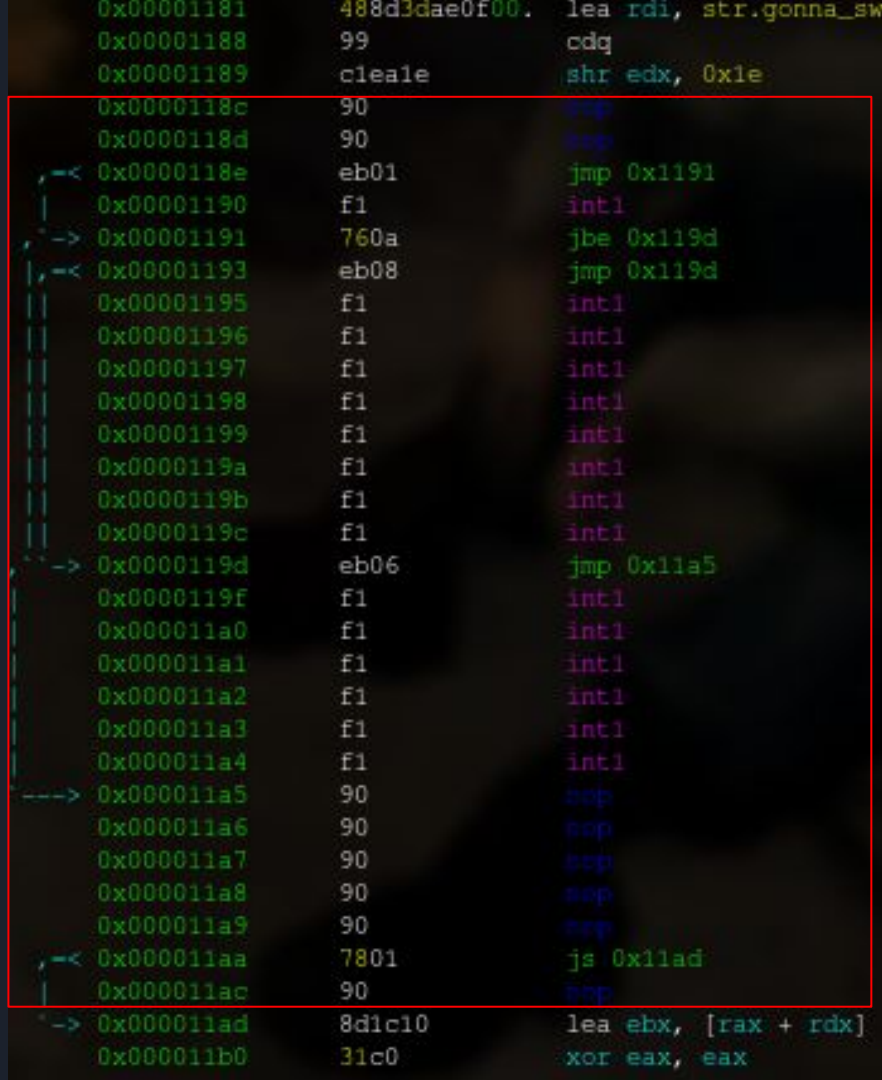

Super Cruftables

Cruftables that can be inserted anywhere, even inside badzones.

Super Cruftables

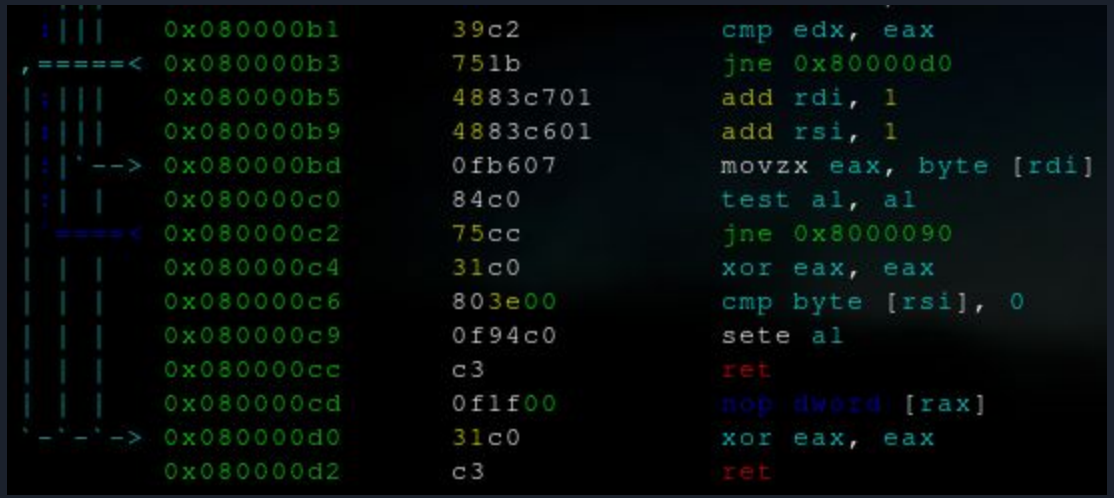

Super Cruftables

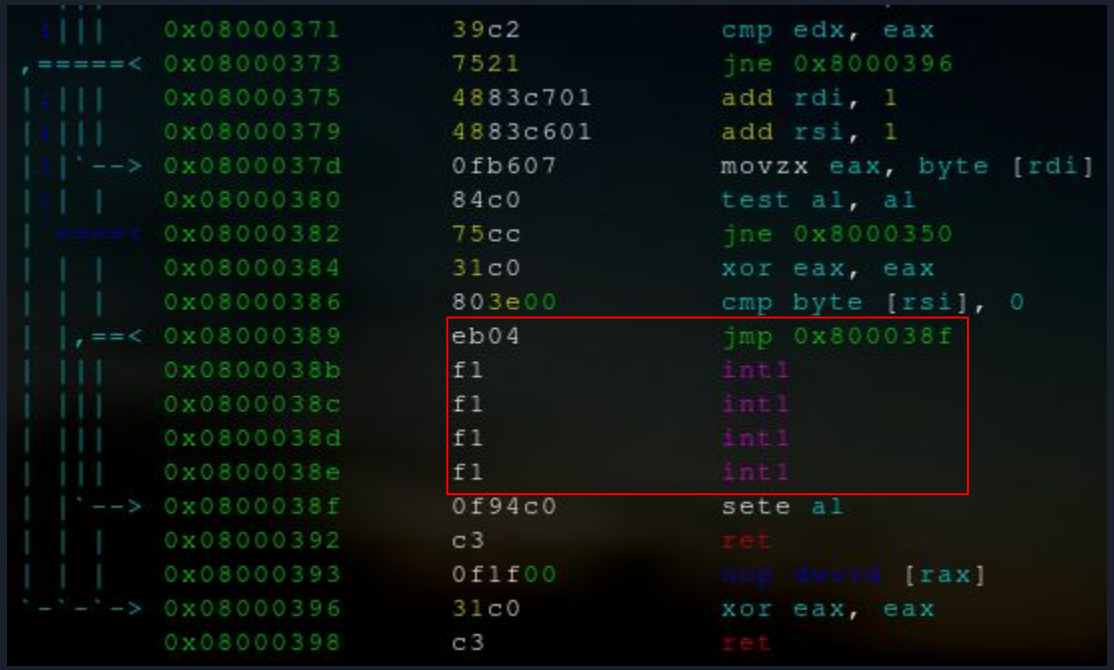

Just update jump offsets

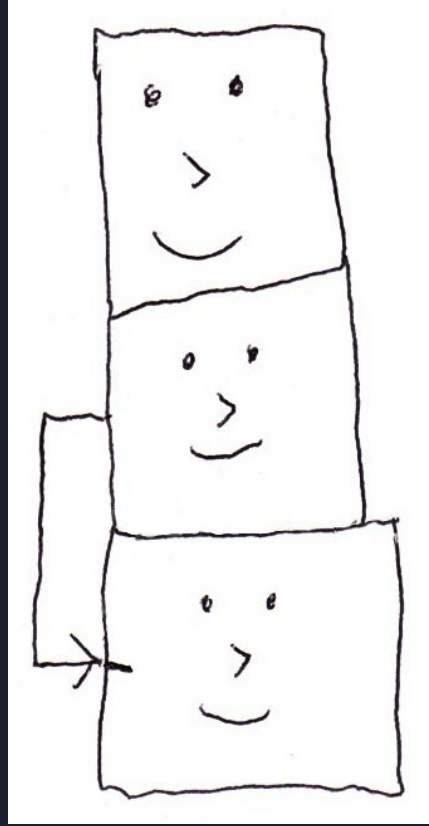

Just update jump offsets

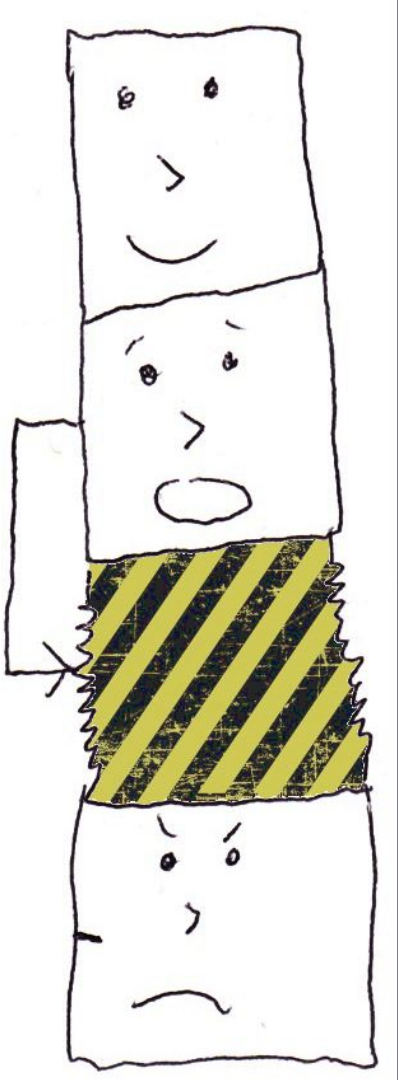

Just update jump offsets

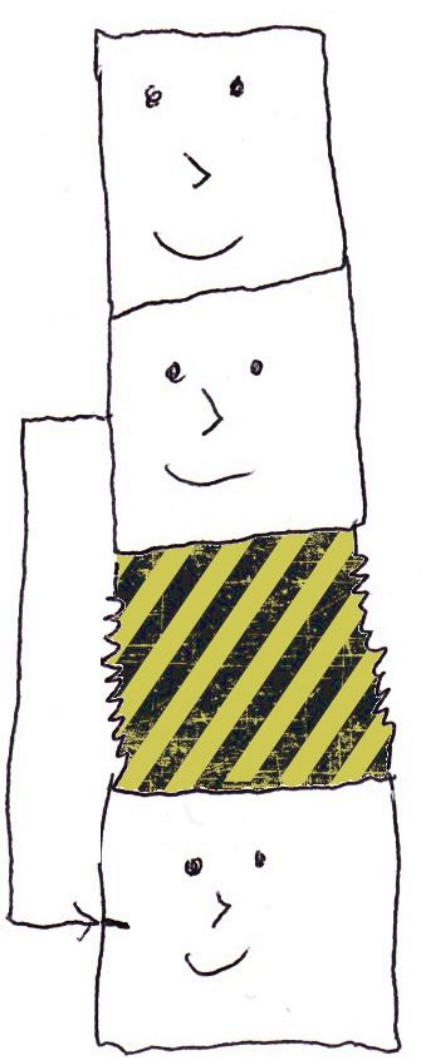

Patch System

- Keep list of all wanted insertions/removals/changes.
  - Apply periodically.
- Attempt each one, ignore if impossible.
- Modifies Jump offsets, symbols, relocs, etc

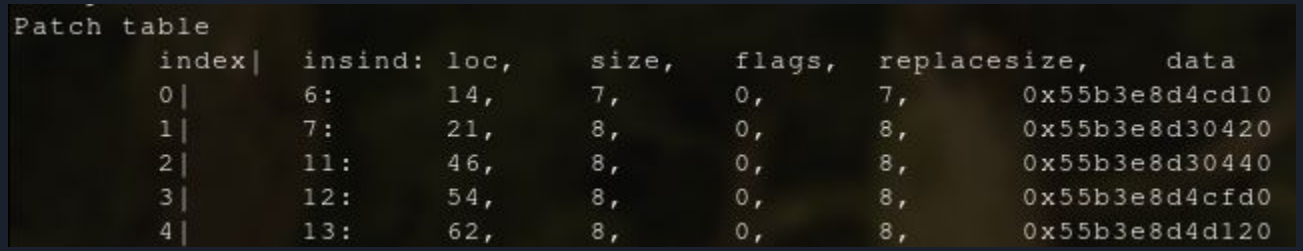

```
Patch table
        index|  insind: loc,    size,   flags,  replacesize,    data
        0|      6:      14,     7,      0,      7,      0x55b3e8d4cd10
        1|      7:      21,     8,      0,      8,      0x55b3e8d30420
        2|      11:     46,     8,      0,      8,      0x55b3e8d30440
        3|      12:     54,     8,      0,      8,      0x55b3e8d4cfd0
        4|      13:     62,     8,      0,      8,      0x55b3e8d4d120
```

Patch System modify jump offsets

- If cruftable is inserted inside a badzone, add cruftable size to jump offset.
  - Instruction start + RAnalOp->nopsize is the start of the offset
  - RAnalOp->size - RAnalOp->nopsize is the size of the offset
    - (int8_t, int16_t, int32_t, etc)

Instruction Bumping

Jumps are Short (1 byte signed) or Near (4 byte signed).

-128 -> 127 vs -2Gigs + 2Gigs

If we go over -128->127 limit, we must “bump” instruction to larger version.

Apply recursively

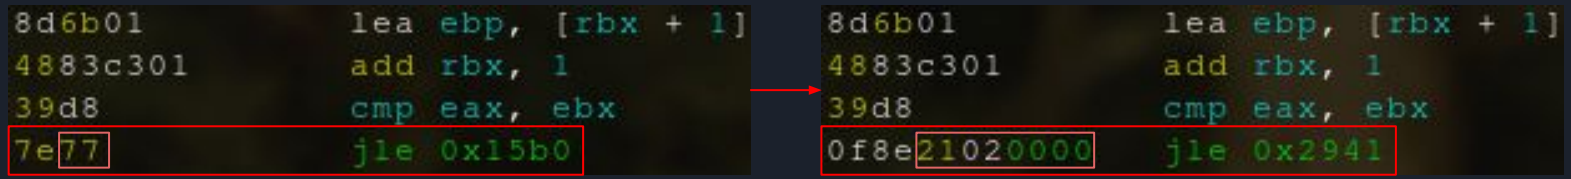

One can’t be bumped

- E3 - JECXZ
  - Only a short jump version, no near jump
- We have to ignore patches that would result in the bumping of JECXZ.
  - Thankfully, rare.

Cruftable obfuscation

Cruftable obfuscation

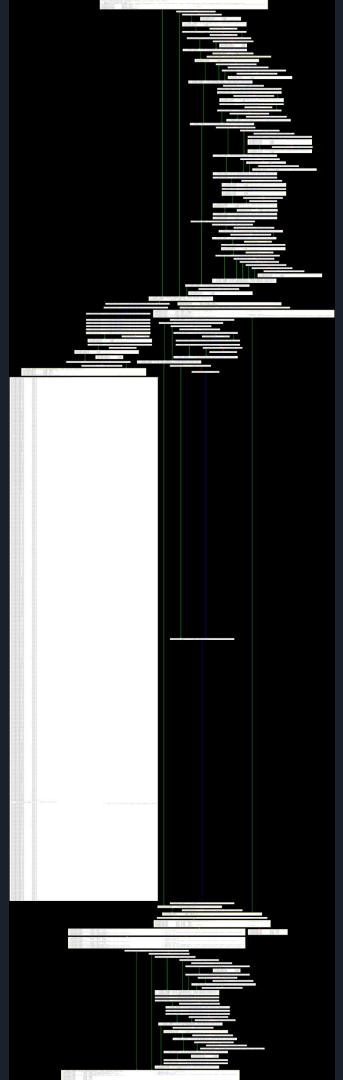

Stack Stuff

- Stack shims
  - Insert a “shim” into each stack frame
    - Changes stack overflow and leak offsets
    - Protects against small overflows
- Stack reordering
  - Reorder Pushes and Pops of saved registers.
    - Changes ROP gadgets.
    - Can change leak offsets.

Stack Shims

- Adjust allocated stack frame size.
  - Add a little “buffer” area on the end.
  - Makes stack-based OOB reads or writes unreliable
  - May even protect against small OOB writes.

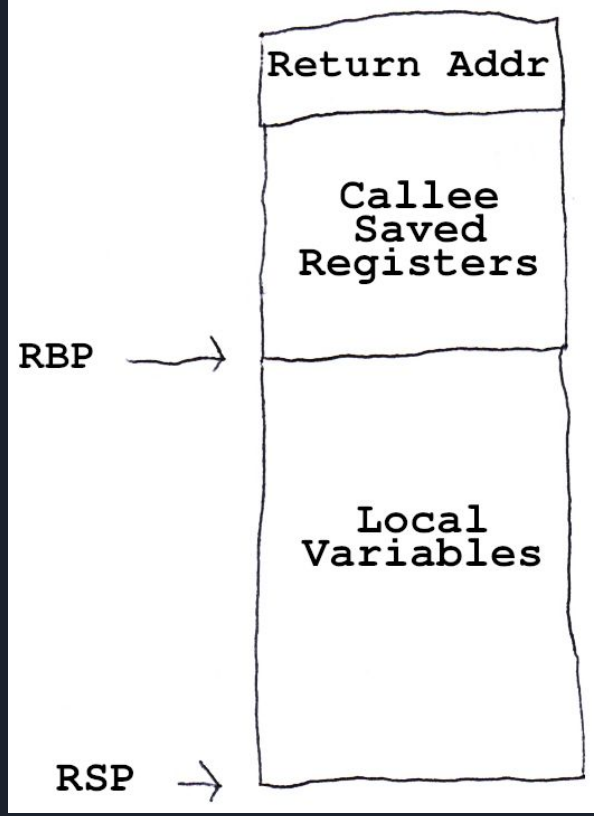

Stack Shims

- Adjust allocated stack frame size.
  - Add a little “buffer” area in frame.
  - Makes stack-based OOB reads or writes unreliable
  - May even protect against small OOB writes.

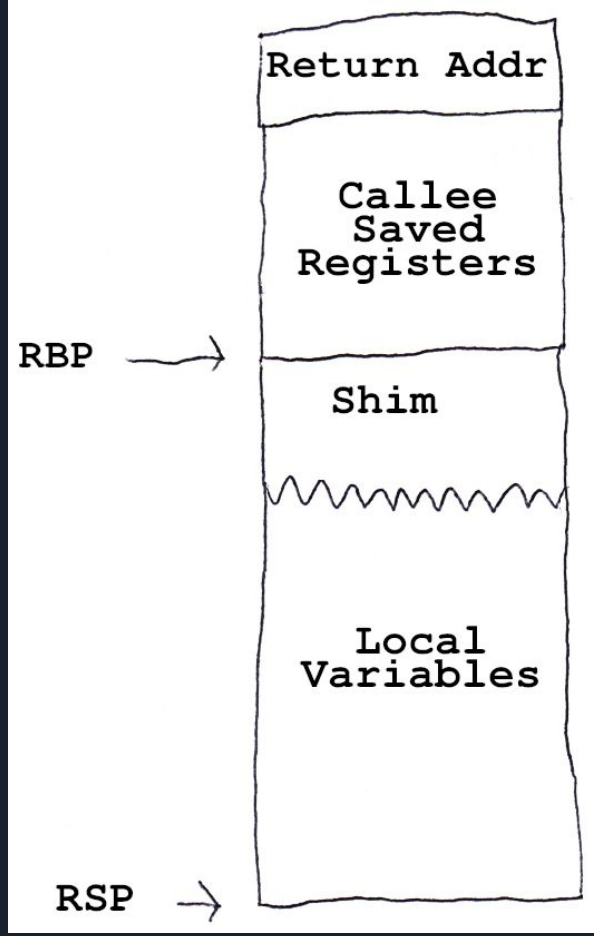

Stack Shim size alignment.

- May have to be aligned to 16 or 8 bytes.
  - GCC defaults to align to 16.
- Some SSE/SIMD instructions will segfault if not aligned.
- Performance impact from unaligned memory accesses.
- **Not always required.**

Stack frame allocation.

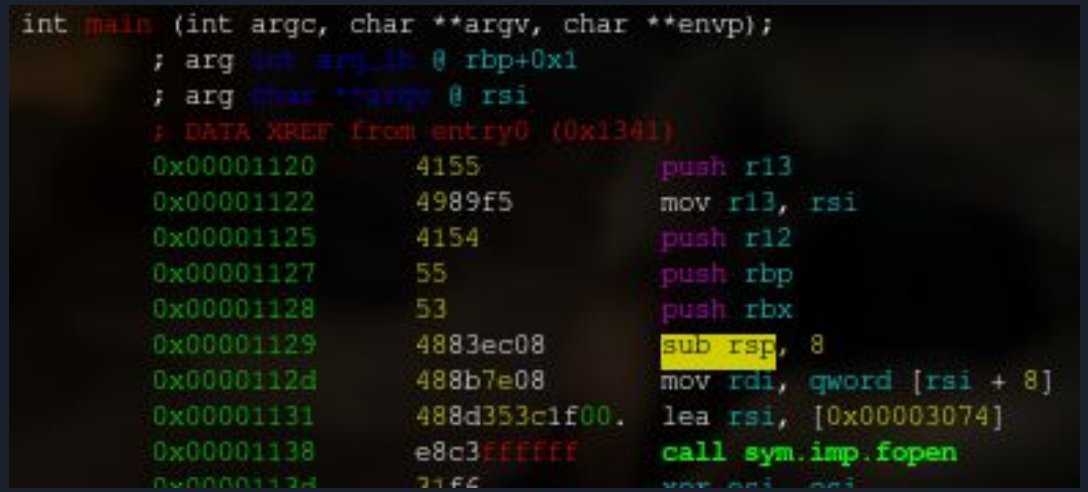

Stack frame deallocation

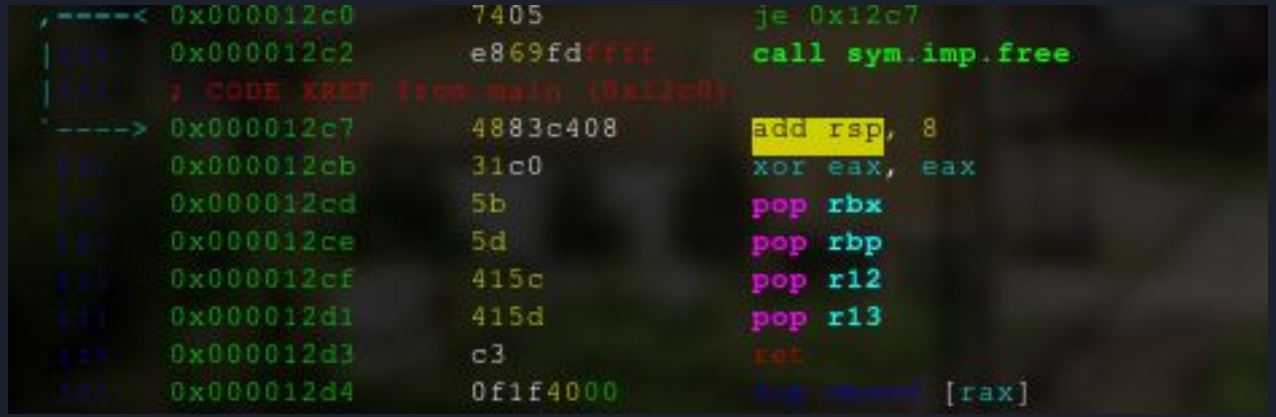

Multiple Deallocations

- Each allocation may have multiple deallocations
  - For allocation in function (symbol), look for all deallocations that match its size.
    - Potential for inconsistencies here, but we've never encountered one in our testing.

Some assumptions

- Only one stack frame per function.
  - Would be difficult to track multiple frames, especially with odd control flow.
  - GCC is nice and does one frame per function (symbol).
- Allocations and deallocations are symmetric
  - No half-allocate or half-deallocate.
  - GCC is also nice here.

Results

```
4883ec48        sub rsp, 0x48
488b15bc2e00.   mov rdx, qword [obj.stdin]

be80000000      mov esi, 0x80
4889e7          mov rdi, rsp
e8affeffff      call sym.imp.fgets
31c0            xor eax, eax
4883c448        add rsp, 0x48
c3              ret
0f1f84000000.   nop dword [rax + rax]
```

```
4883ec68        sub rsp, 0x68
488b15bc2e00.   mov rdx, qword [obj.stdin]

be80000000      mov esi, 0x80
4889e7          mov rdi, rsp
e8affeffff      call sym.imp.fgets
31c0            xor eax, eax
4883c468        add rsp, 0x68
c3              ret
0f1f84000000.   nop dword [rax + rax]
```

Stack Reordering

Reorder the pushes and pops of saved registers.

Two places to reorder

- Prologue
  - Only one, the start of the function. (Symbol)
  - Lots of Pushes

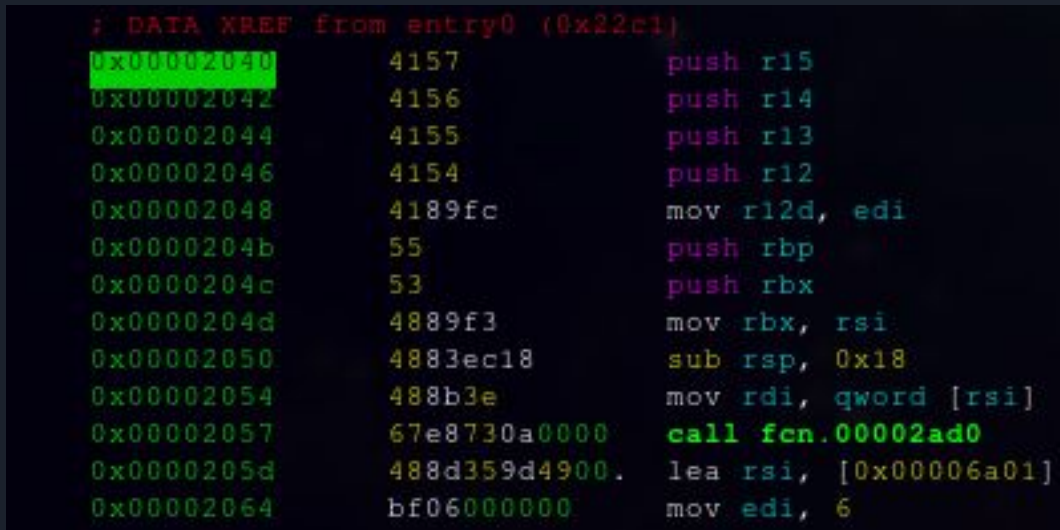

Two places to reorder

- Epilogue(s)
  - Can be multiple
  - Right before a “ret”
  - Lots of Pops
  - May not be directly at the end of a function.

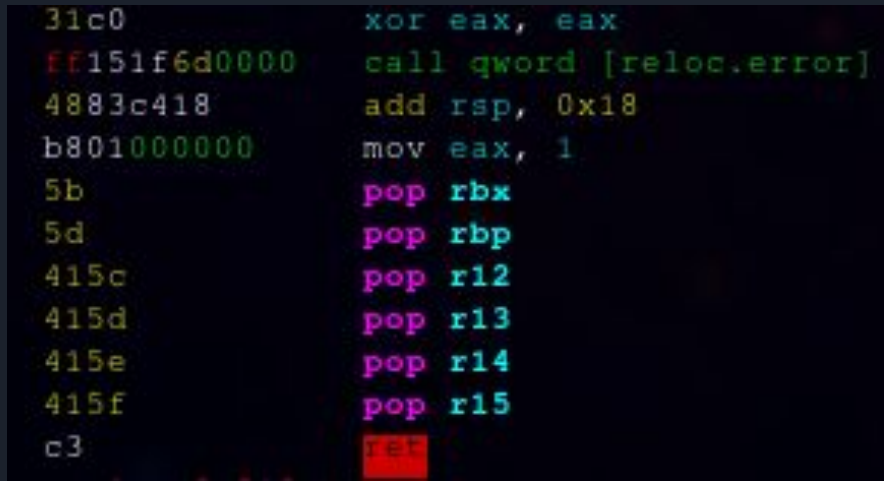

Look for reorderable instructions

- Requirements:
  - **No modifications to RIP** (Jumps, calls, etc)
  - No jumps pointing into the area (Look at badzones).
  - No relocs in the area (may be able to be changed in the future)
  - **No stack operations that aren't push/pop.**
    - **strstr(op->esil, 'rsp')**
  - **No memory Writes** (might be not required)
  - **Can not use a push if in a epilogue, can not use a pop if in an Prologue.**
  - Register pushes must match Register Pops.
    - Truncation step.

RAnalOp Is Super Useful. (Again)

```
typedef struct r_anal_op_t {
        char *mnemonic; /* mnemonic.. it actually contains the args too. we should replace rasm with this */
        ut64 addr;      /* address */
        ut32 type;      /* type of opcode */
        ut64 prefix;    /* type of opcode prefix (rep,lock,..) */
        ut32 type2;     /* used by java */
        int group;      /* is fpu, is privileged, mmx, etc */
        int stackop;    /* operation on stack? */
        int cond;       /* condition type */
        int size;       /* size in bytes of opcode */
        int nopcode;    /* number of bytes representing the opcode (not the arguments) TODO: find better name */
        int cycles;     /* cpu-cycles taken by instruction */
        int failcycles; /* conditional cpu-cycles */
        int family;     /* family of opcode */
        int id;         /* instruction id */
        bool eob;       /* end of block (boolean) */
        bool sign;      /* operates on signed values, false by default */
        /* Run N instructions before executing the current one */
        int delay;      /* delay N slots (mips, ..)*/
        ut64 jump;      /* true jmp */
        ut64 fail;      /* false jmp */
        int direction;  /* 1 = read, 2 = write, 4 = exec, 8 = reference,  */
        st64 ptr;       /* reference to memory */ /* XXX signed? */
        ut64 val;       /* reference to value */ /* XXX signed? */
        int ptrsize;    /* f.ex: zero extends for 8, 16 or 32 bits only */
        st64 stackptr;  /* stack pointer */
        int refptr;     /* if (0) ptr = "reference" else ptr = "load memory of refptr bytes" */
        RAnalVar *var;  /* local var/arg used by this instruction */
        RAnalValue *src[3];
        RAnalValue *dst;
        struct r_anal_op_t *next; // TODO deprecate
        RStrBuf esil;
        RStrBuf opex;
        const char *reg; /* destination register */
        const char *ireg; /* register used for indirect memory computation*/
        int scale;
        ut64 disp;
        RAnalSwitchOp *switch_op;
        RAnalHint hint;
        RAnalDataType datatype;
} RAnalOp;
```

Truncation step

- If a Prologue/epilogue has an extra push or pop that isn't matched in ALL of the others, we must truncate it out.
  - Rare, but can happen.

What do we have left?

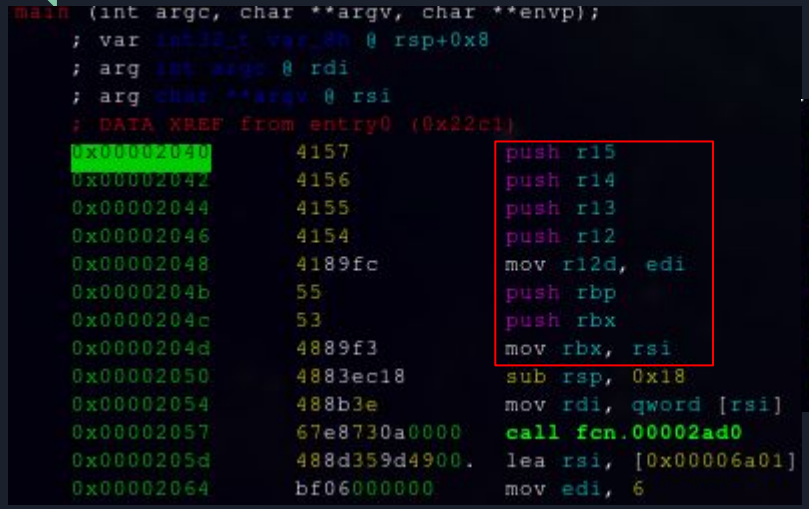

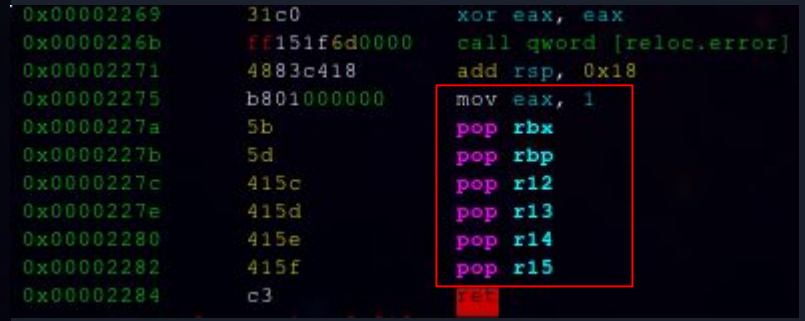

Build dependency list

- Each instruction has a list of dependency instructions that must come before.
  - For every instruction, see what instructions comes before it that use any of the same registers.
  - Instructions may have many dependencies, but cycles are not possible.
- When shuffling, instruction must have all of its dependencies fulfilled before being chosen.
  - (Fisher-yates)
- Seperate lists for every epilogue/prologue.

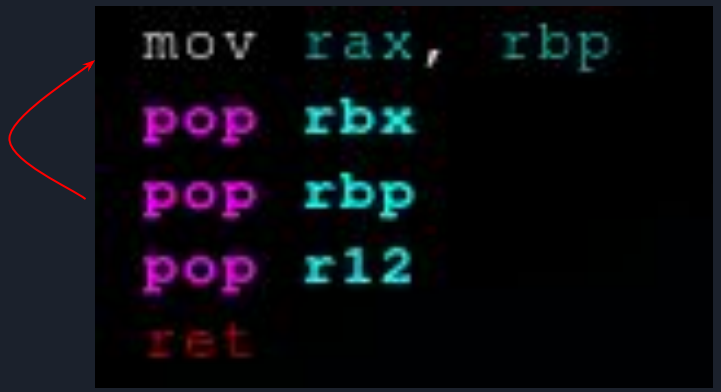

Build dependency list

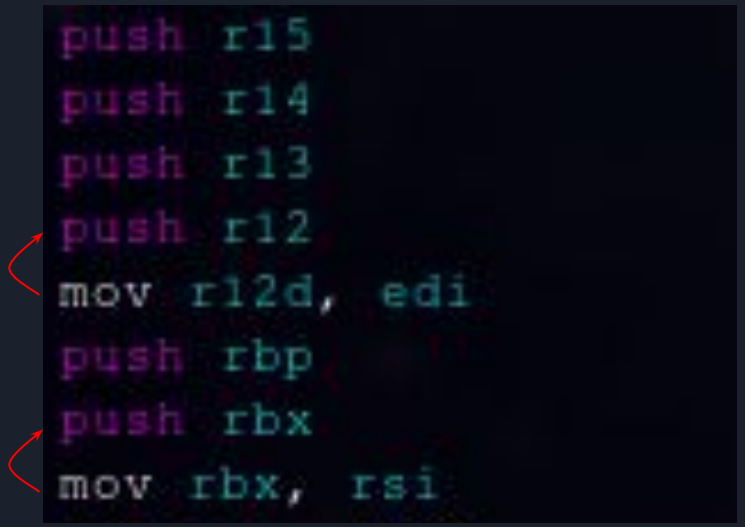

Dependencies kept after shuffling

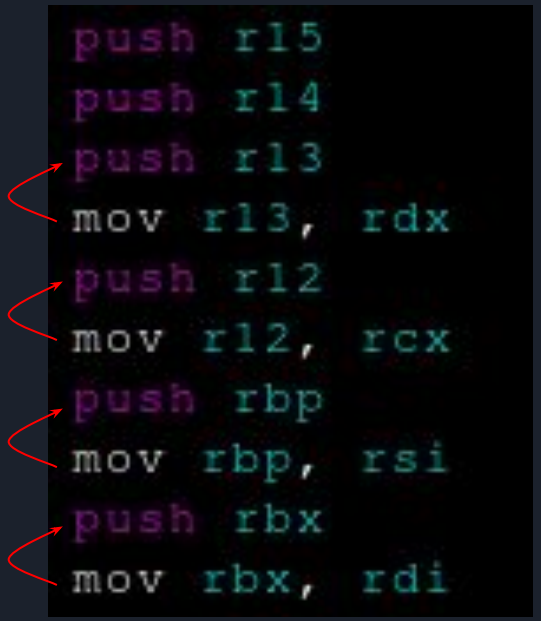

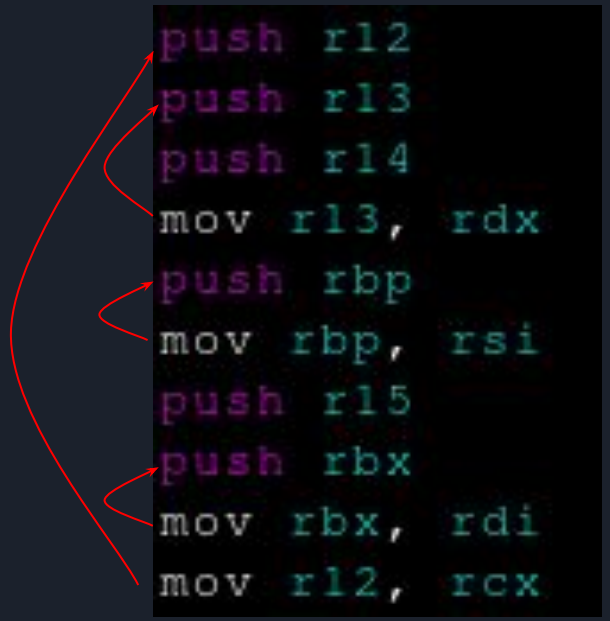

Prologue/Epilogue Symmetry

- Whatever order we shuffle the Pushes in the prologue, we must change the Pops in Epilogue to match.
  - Reverse order.
  - Only pushes/pops need to follow order. Misc instructions are fine to reorder.
  - Still have to follow instruction dependencies too.

Prologue/Epilogue Symmetry

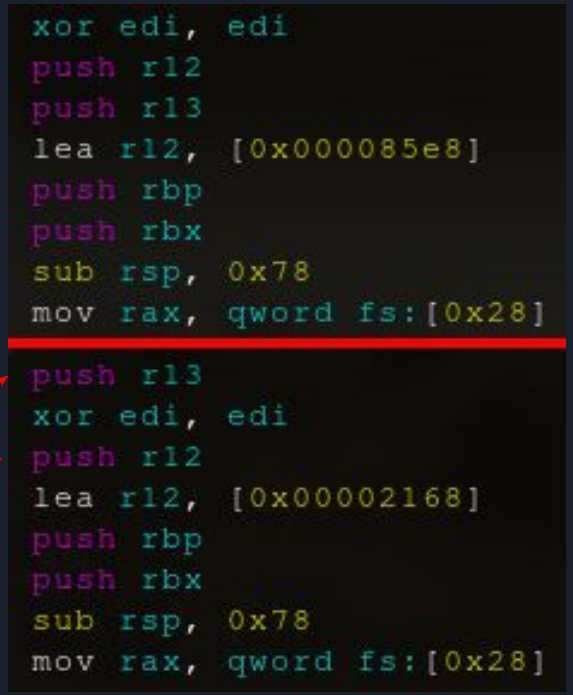

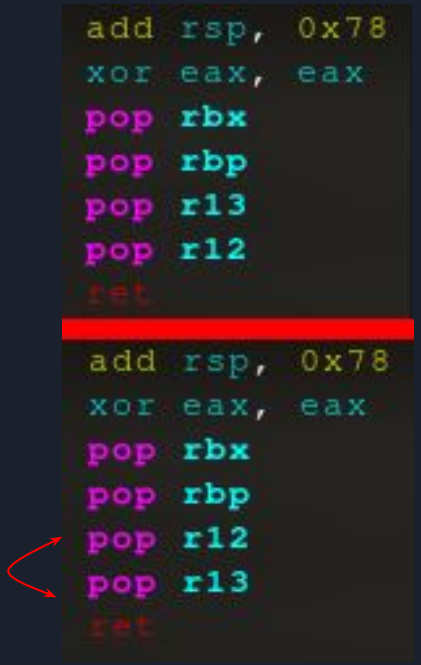

Demo

Future things

- Use R2 emulation functionality for fancier analysis/verification.
  - Would allow for more complex modifications and higher reliability.
- Port to ARM, Mips, Risc-V, etc.
- Full binary support.
  - Might be difficult due to linker resolving some symbols.
- Port to PE, Mach-O, etc.
- More advanced cruftable data.
  - Chunks of code that actually look like code.
  - More complex structure.
  - Harder to automatically de-obfuscate.
- Shuffle at binary/object load time
  - Kernel module?
  - It's fast enough for this to be feasible.
- Standalone "Patch System" as a plugin for R2.

Special thanks,

Martin Osterloh for section shuffling.

Jesse Earisman for most of the ELF code.

R2 devs for answering all of my noob questions and fixing bugs.

Siege Technologies for sending me to R2Con.

Questions?

Unused/unfinished slides

Jumptable hack

GCC is “smart” and does jumptables in a weird way.

Todo

Instruction Bumping

- Convert short jumps to near jumps. (if needed)
- Processed recursively by patch system.
  - Bumping one instruction may cause another instruction to be bumped.

```
BUMPING -121, -1337: 1 -> 2
```

Analysis output

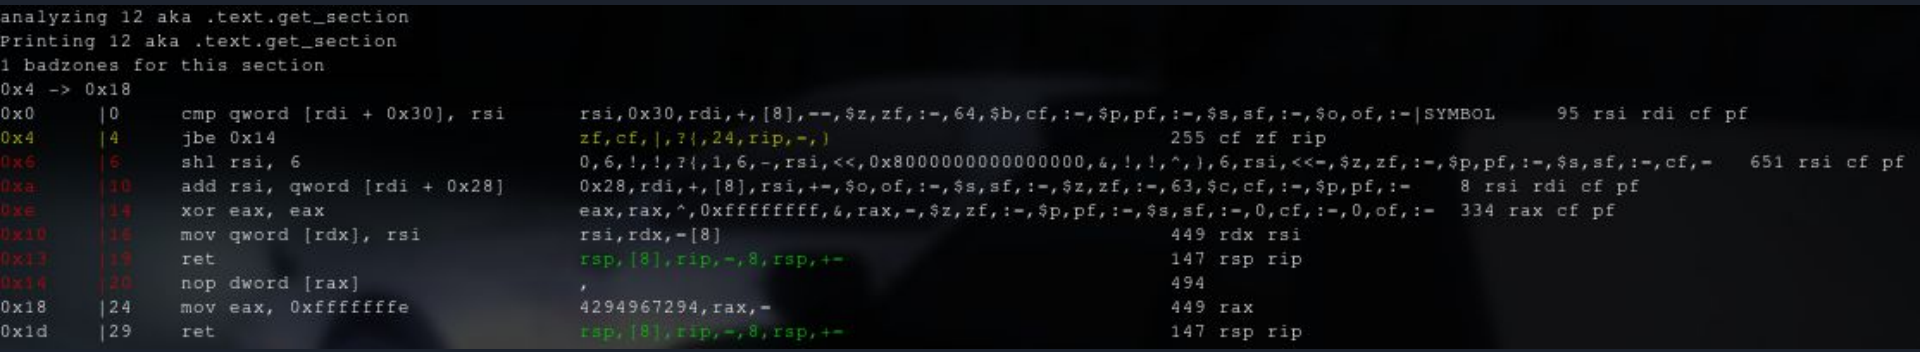

```
analyzing 12 aka .text.get_section
Printing 12 aka .text.get_section
1 badzones for this section
0x4 -> 0x18
0x0     |0     cmp qword [rdi + 0x30], rsi      rsi,0x30,rdi,+,[8],--,$z,zf,:-,64,$b,cf,:-,$p,pf,:-,$s,sf,:-,$o,of,:-|SYMBOL     95 rsi rdi cf pf
0x4     |4     jbe 0x14                         zf,cf,|,?{,24,rip,-,}                                                  255 cf zf rip
0x6     |6     shl rsi, 6                       0,6,!,!,?{,1,6,-,rsi,<<,0x8000000000000000,&,!,!,^,},6,rsi,<<-,$z,zf,:-,$p,pf,:-,$s,sf,:-,cf,-   651 rsi cf pf
0xa     |10    add rsi, qword [rdi + 0x28]      0x28,rdi,+,[8],rsi,+=,$o,of,:-,$s,sf,:-,$z,zf,:-,63,$c,cf,:-,$p,pf,:-    8 rsi rdi cf pf
0xe     |14    xor eax, eax                     eax,rax,^,0xffffffff,&,rax,-,$z,zf,:-,$p,pf,:-,$s,sf,:-,0,cf,:-,0,of,:-  334 rax cf pf
0x10    |16    mov qword [rdx], rsi             rsi,rdx,=[8]                                                           449 rdx rsi
0x13    |19    ret                              rsp,[8],rip,-,8,rsp,+=                                                 147 rsp rip
0x14    |20    nop dword [rax]                  ,                                                                      494
0x18    |24    mov eax, 0xfffffffe              4294967294,rax,-                                                       449 rax
0x1d    |29    ret                              rsp,[8],rip,-,8,rsp,+=                                                 147 rsp rip
```

Two types of relative addressing

- RSP and RBP relative.
  - RSP+offset
    - for local variables.
  - RBP+offset
    - for stack arguments.
- (Ideally)

GCC sometimes does gross stuff

And this changes from version to version.

- RSP + to grab stack arguments.
  - “Reaches over rbp”
- RBP- relative addressing for local vars.

Addressing Fixes

- Convert the “bad” modes to the appropriate one. - Will not work in all cases.
  - If RBP is used as general purpose, can’t convert to RBP-relative.
  - May need to do some tracking of pushes/pops between frame allocation and usage.
- Add the stack shim offset into the bad modes. - More likely to work.
  - Needs to be done to every instruction that “Reaches over”

We currently use the first method. Not 100% reliable.

GCC Sibling calls

- GCC for optimization will make “sibling calls”
  - Tail-call optimization
  - Instead of a call/ret, just use a JMP.
- Or sometimes will split a function in two.
- Force to show up as a “epilogues”.
  - There is no ret.
  - May not fully pop stuff from stack. Symmetry/truncation will handle it.